

Enterprise Knowledge Based Automation

Complete Architecture & Operations Handbook

Tumhare apne project ka hands-on training manual — code, Terraform, scripts aur AWS deployment ke hisaab se, simple Hinglish mein.

Repository: [Analyst-Parveen/Enterprise-Knowledge-Based-Automation](#)
Snapshot: 2026-09-12 · branch [aws-deployment](#) · tenant onboarding ke baad

Contents

1. Is Handbook Ko Kaise Use Karein
2. Project Overview — Ye Project Karta Kya Hai?
 - 2.1 Problem kya hai?
 - 2.2 Kaun use karta hai?
 - 2.3 Important terms — simple bhasha mein
 - 2.4 Real-world example
3. Complete Architecture
 - 3.1 Poora system ek nazar mein
 - 3.2 Kya live hai, kya nahi — honest table
4. Project Structure — Kaunsi File Kya Karti Hai?
 - 4.1 Repository tree (important hisse)
 - 4.2 Important files — ek-ek karke
5. AWS Services — Ye Service Kya Kar Rahi Hai?
 - 5.1 Terraform — sab kuch banane wala
 - 5.2 S3 — original documents
 - 5.3 Cognito — user identity
 - 5.4 ECS Fargate — backend chalane wali machine
 - 5.5 ECR — Docker images ka godown
 - 5.6 ALB — backend ka darwaaza
 - 5.7 CloudFront — API ka HTTPS front door
 - 5.8 Amplify — frontend hosting
 - 5.9 CodeDeploy — blue/green release
 - 5.10 Secrets Manager — secrets
 - 5.11 CloudWatch — logs aur alarms
 - 5.12 Bedrock — AI models
 - 5.13 Amazon Transcribe
 - 5.14 Lambda, SNS, EventBridge, Budgets — cost guard
 - 5.15 IAM — kaun kya kar sakta hai
 - 5.16 VPC & security groups
 - 5.17 RDS — PostgreSQL database (managed)
 - 5.18 Jo services is project mein NAHI hain
6. Authentication — Login Kaise Kaam Karta Hai
 - 6.1 Flow — ek nazar mein
 - 6.2 Basic terms
 - 6.3 Frontend kya karta hai (actual code)
 - 6.4 Backend kya karta hai (core/auth.py)
 - 6.5 Tenant aur role kahan se aate hain
 - 6.6 ID token vs Access token — kyun ID token
 - 6.7 Password kabhi app ke paas nahi aata
 - 6.8 Token expiry
7. Company / Tenant / User Creation
 - 7.1 Seedhi baat pehle
 - 7.2 Tenant ID kahan "rehta" hai?
 - 7.3 Example: "ABC Insurance" ko onboard karna
 - 7.4 Local mein users (no Cognito)
8. Admin vs User
 - 8.1 Kaun kya kar sakta hai (code ke hisaab se)
 - 8.2 Limits (teeno roles par)
 - 8.3 Scenario
9. Frontend — Browser Mein Kya Chal Raha Hai
 - 9.1 Tech stack (frontend/package.json)
 - 9.2 Do build modes
 - 9.3 Pages (routes)
 - 9.4 API client — src/lib/api.ts
 - 9.5 Loading / empty / error states
 - 9.6 Frontend → backend path
10. Backend — FastAPI Andar Se
 - 10.1 Request ka safar (har request)
 - 10.2 Endpoints (sab /api/v1 ke neeche)
 - 10.3 Error format
11. Database — PostgreSQL
 - 11.1 Technology
 - 11.2 Tables (10)
 - 11.3 ER diagram (simplified)
 - 11.4 Indexes (tenant filter fast rakhne ke liye)
 - 11.5 Migrations
12. Document Ingestion Flow
 - 12.1 Diagram
 - 12.2 Har stage

- 12.3 Status kaise badalta hai
- 12.4 Manually check karna
- 13. RAG Query Flow — Sabse Important Chapter
 - 13.1 Example
 - 13.2 Har step — Hinglish mein
 - 13.3 Response envelope (frontend ko kya milta hai)
 - 13.4 Hallucination kaise kam hota hai
 - 13.5 Local mode vs AWS mode
 - 13.6 Agentic workflows (LangGraph)
- 14. AI / LLM / Embeddings
 - 14.1 Models
 - 14.2 Pricing table (code mein, ai/registry.py, per 1M tokens)
 - 14.3 Flow
- 15. Qdrant — Vector Database
- 16. Redis — Cache aur Rate Limit
- 17. Data Storage Map — Kya Kahan Rehta Hai
- 18. Local Development — Laptop Par Poora Project
 - 18.1 Picture
 - 18.2 Sabse aasaan raasta — ek command
 - 18.3 Haath se, step by step (RUNBOOK Part A se verified)
 - 18.4 Check, logs, docs
 - 18.5 Stop / restart
- 19. AWS Production Flow
 - 19.1 Frontend
 - 19.2 Backend
 - 19.3 Blue/green — kaise kaam karta hai
- 20. "Mujhe Code Change Karna Hai — Ab Kya Karu?"
 - 20.1 Backend change
 - 20.2 Frontend change
 - 20.3 Decision tree
- 21. Branch Strategy — Asli Setup
- 22. Stop / Start / Destroy — Cheat Sheet
 - 22.1 Stop vs Destroy
 - 22.2 Commands
 - 22.3 destroy.sh kya karta hai
 - 22.4 Destroy ke baad redeploy
 - 22.5 Rollback — galat version live ho gaya to kya karu?
- 23. Docker
- 24. Terraform — Beginner Guide
 - 24.1 Basic commands
 - 24.2 Chaar states
 - 24.3 "Plan mein destroy/replace dikha — ab kya karu?"
 - 24.4 State lock atka ho
- 25. CI/CD — GitHub Actions
 - 25.1 ci.yml (automatic)
 - 25.2 deploy.yml (manual, workflow_dispatch)
- 26. Testing
- 27. Monitoring & Logs — "Problem Aaya To Sabse Pehle Kahan Dekhu?"
- 28. Troubleshooting Playbook
 - 28.1 "Kuch toot gaya" — decision tree
 - 28.2 Problem → Check → Command → Wajah → Fix
- 29. Security
 - 29.1 Controls — kya kahan
 - 29.2 Honest gaps
 - 29.3 GitHub mein kabhi commit nahi karna (.gitignore se)
- 30. Cost Guard — Kharcha Kaise Control Hota Hai
 - 30.1 Pehle ek sach
 - 30.2 Architecture
 - 30.3 Thresholds
 - 30.4 Shutdown kya karta hai / kya chhodta hai
 - 30.5 DRY_RUN
 - 30.6 Manually verify
 - 30.7 Kharcha — numbers
- 31. Manual Data Inspection (read-only)
- 32. Real End-to-End Example — "ABC Insurance"
- 33. Complete Command Cheat Sheet
- 34. Interview Preparation
 - 34.1 Two-minute version
 - 34.2 Five-minute version — structure
 - 34.3 Likely questions — short answers

- 35. NEVER DO THIS
- 36. Glossary
- 37. Final Quick Reference

1. Is Handbook Ko Kaise Use Karein

Parveen, ye handbook tumhare **apne repository** ko padh kar likhi gayi hai — koi generic AWS ya RAG tutorial nahi. Har command, har file path, har AWS resource ka naam code, Terraform aur scripts se verify kiya gaya hai.

Teen rules jo is handbook mein follow kiye gaye hain:

Rule	Matlab
Code hi source of truth hai	Agar docs aur code alag baat kehte hain, to dono likhe gaye hain aur bataya gaya hai ki code kya karta hai
"Implemented nahi hai" honestly likha hai	Jo cheez repo mein nahi hai, uske liye likha hai: Project mein currently ye implemented nahi hai.
Koi secret nahi	Passwords, tokens, keys, account ID — kuch nahi. Unki jagah <code><YOUR_...></code> placeholders hain

Kaise padhein:

- Pehli baar: Chapter 2 se 5 tak seedha padh lo — project samajh aa jayega.
- Kuch karna hai (deploy, change, stop): Chapter 18–24 aur 33 (cheat sheet).
- Kuch toot gaya: Chapter 28 (troubleshooting) aur Chapter 27 (logs).
- Interview se pehle: Chapter 34.

Box ke rang ka matlab: neela = samjhane wali baat · peela = dhyaan do · laal = ye kabhi mat karna · hara = verified fact / achhi practice.

Is snapshot ki live state (2026-09-12):

Cheez	State
AWS account plan	Paid (Free plan CodeDeploy block karta hai)
Region	<code>us-west-2</code>
Backend stack (<code>envs/dev</code>)	Deployed — ECS + ALB + CodeDeploy blue/green
Database	Amazon RDS PostgreSQL 16 (<code>ekba-dev-postgres</code> , <code>db.t4g.micro</code> , private) — Postgres sidecar ki jagah. Task replace hone par data bachta hai (2026-09-12 ko row counts se verify kiya)
Frontend (Amplify)	Live: <code>https://aws-deployment.d39pgpamq0p0n8.amplifyapp.com</code>
API front door (CloudFront)	Live: <code>https://d2jw2wchz5oekh.cloudfront.net</code>
Cost guard	Applied, <code>DRY_RUN = true</code> (report karta hai, kuch rokta nahi). Ab RDS ko stop bhi karta hai (delete kabhi nahi)
Bedrock quota	0 — isliye AWS par chat answer aur demo seeding abhi nahi chalte
Roles	Teen: <code>user</code> , <code>admin</code> , <code>platform_admin</code> (Chapter 2.2)
Tenant onboarding	Implemented — platform operator company banata hai aur uska pehla admin invite karta hai; wo admin apne users banata hai (Chapter 7)
Login	Email + password (Cognito), pehli baar password set karna, forgot-password, silent refresh, server-side logout. Token paste karna sirf local debugging ke liye bacha hai (Chapter 6)
Tests	280 backend (deploy gate 279) + 29 Playwright journeys

Ye URLs secret nahi hain, lekin ye state badal sakti hai. `destroy.sh` ke baad backend ALB ka naam badal jata hai; CloudFront aur Amplify URL same rehte hain (Chapter 22).

2. Project Overview — Ye Project Karta Kya Hai?

2.1 Problem kya hai?

Ek company ke paas bahut saare documents hote hain — HR policy, travel policy, NDA template, sales deck, safety SOP, spreadsheets, diagrams, recordings. Employee ko ek chhota sa sawaal poochna ho ("hotel reimbursement limit kitni hai?") to use pooraa PDF dhoondhna padta hai.

Ye project ek **private company knowledge assistant** hai:

- Company apne documents upload karti hai.
- System unhe padh kar, chhote tukdon (chunks) mein tod kar, searchable bana deta hai.
- Employee natural language mein sawaal poochta hai.
- System **sirf usi company ke** documents mein se relevant hissa dhoondh kar, AI model se answer banwata hai — **source (citation) ke saath**.

2.2 Kaun use karta hai?

Role	Kaun	Code mein kahan
user	Normal employee	app/core/context.py — Role = "user" "admin" "platform_admin"
admin	Us company ka admin	Same; admin endpoints app/api/v1/admin.py
platform_admin	Service provider khud — jo companies ko onboard karta hai	Same; platform endpoints app/api/v1/platform.py

Sirf yahi **teen** roles hain, aur code inke alawa kuch accept nahi karta — token mein koi chautha role aaya to request 401 hoti hai, downgrade nahi.

`platform_admin` ek reserved tenant `platform` mein rehta hai. Wo tenant sirf operators rakhta hai — na documents, na conversations. Isliye tenant filtering mein koi exception banane ki zaroorat nahi padti: platform role bhi apne hi tenant se bandha hai.

Hierarchy — kaun kya kar sakta hai:

Role	Kar sakta hai	Kabhi nahi kar sakta
platform_admin	company banana, uska pehla admin invite karna, onboarding trail padhna	kisi bhi company ke documents, chat, ya metrics dekhna
admin	sirf apni company ke users banana/manage karna, user ya admin role dena, apni company ke metrics aur audit log	nayi company banana, platform_admin banana, doosri company ko chhoona
user	product use karna	users ya companies manage karna

Teen cheezein is hierarchy ko todne se rokti hain:

1. **Platform role aur platform tenant ek doosre ko imply karte hain.** Token verify karte waqt hi check hota hai — `platform_admin` normal tenant mein, ya normal role platform tenant mein, dono reject. Aadha identity forge karna kaam nahi karta.
2. **Koi API platform_admin nahi de sakti.** Request schema mein hi ye value allowed nahi (422 milta hai handler chalne se pehle), aur server-side `assert_role_assignable` dobara refuse karta hai. Pehla operator `scripts/bootstrap-platform-admin.sh` se banta hai — AWS credentials wale insaan ke haath se, product ke bahar.
3. **Cross-tenant read sirf ek jagah hai, aur audited hai.** `app/db/control_plane.py` — sirf registry data (naam, id, seat counts), `PlatformAdminUser` dependency ke peeche, har mutation ka audit event.

2.3 Important terms – simple bhasha mein

Multi-tenant – Ek hi application, kai companies use karenin, aur har company ka data doosri se bilkul alag rahe. "Tenant" matlab ek company. Is project mein har company ka ek `tenant_id` hota hai (jaise `seed-tenant-northwind`).

RAG (Retrieval-Augmented Generation) – AI model se seedha answer poochne ke bajaye, pehle apne documents mein se relevant tukde **dhoondho (retrieve)**, phir wo tukde model ko do aur kaho "sirf inke basis par answer do". Isse hallucination (banayi hui baat) kam hoti hai.

Agentic RAG – Ek se zyada step wala workflow jahan system khud plan karta hai, retrieve karta hai, analyze karta hai, phir answer banata hai. Ye project mein **implemented hai**: `LangGraph graph plan → retrieve → analyze → synthesize → cite` (`app/services/agents/graph.py`), 5 workflows ke saath (Chapter 13.6).

Citations – Answer ke saath ye batana ki ye baat kis document ke kis hisse se aayi. Model answer mein `[S1]`, `[S2]` jaise markers lagata hai; backend check karta hai ki wo marker sach mein retrieve kiye gaye chunk ka hai (`app/services/rag/guardrails.py → validate_citations`).

Guardrails – Suraksha ki deewarein. Input ke liye (prompt injection check), aur output ke liye (system prompt leak, secrets, unsafe HTML). Is project mein ye pipeline ke **compulsory stages** hain, optional nahi.

Tenant isolation – Company A ka user kabhi bhi Company B ka ek bhi document, chunk, cache entry ya answer na dekh sake. Ye is project ka **sabse important rule** hai.

2.4 Real-world example


Verified: ye exact example seed data mein hai — `python -m seeds.seed "Travel Policy 2026.pdf"` banata hai aur `RUNBOOK` ke hisaab se answer "150 USD" aata hai (local mode mein).

3. Complete Architecture


3.1 Poora system ek nazar mein

Neeche ke diagrams sirf wahi components dikhate hain jo repository mein actually configure hain. Kaunsa component live hai aur kaunsa sirf configured / blocked (jaise Bedrock quota 0), wo section 3.2 ki table mein hai.


(A) Request path – browser se containers tak:



(B) api container kin AWS services se baat karta hai:



Cost guard aur baaki control plane (alag diagram - Chapter 30):



3.2 Kya live hai, kya nahi – honest table

Component	Status	Detail
Amplify frontend	Live	Static export, aws-deployment branch se auto-build
CloudFront API	Live	HTTPS layer, ALB HTTP-only hai isliye
ALB + ECS Fargate	Live	Ek task, teen containers (api, qdrant, redis)
PostgreSQL	Live (Amazon RDS)	ekba-dev-postgres, db.t4g.micro, private subnets. Data task restart par bachta hai ; destroy.sh snapshot leta hai
Supabase	Implemented nahi hai	Project mein currently ye implemented nahi hai
Qdrant	Live (container)	Sidecar, koi managed vector DB nahi
Redis	Live (container)	Sidecar. ElastiCache nahi
S3	Live	Documents bucket (baseline state)
Cognito	Live (code path)	Pool + client, aur poora email/password sign-in + invitation flow implemented; AWS par live user login pending hai (Chapter 6)
Bedrock	Configured, blocked	Account quotas 0 – chat/embedding AWS par throttle hote hain
Transcribe	Configured	Audio/video ingestion ke liye; AWS par test nahi hua
CodeDeploy	Live	Blue/green deployments successful
CloudWatch Logs	Live	/ekba/dev/service (1 din retention)
CloudWatch Alarms	Live	ekba-dev-5xx, ekba-dev-latency
Cost guard (Budgets, SNS, Lambda, EventBridge)	Live, DRY_RUN	Report karta hai, act nahi karta
LangSmith	Configured only	Setting + secret hai, lekin tracing ka code implemented nahi hai
SQS	Implemented nahi hai	Project mein currently ye implemented nahi hai
EKS / NAT Gateway / ElastiCache / EFS	Kabhi nahi	Cost rule ke hisaab se Terraform mein allowed hi nahi
RDS Multi-AZ / bada instance / doosra database	Kabhi nahi	Sirf ek db.t4g.micro allowed hai (Chapter 11)

4. Project Structure – Kaunsi File Kya Karti Hai?

4.1 Repository tree (important hisse)

4.2 Important files – ek-ek karke

Har file ke liye: Purpose → Kaun call karta hai → Ye kya call karti hai → Change ka asar → Kab edit karna → Kaise test karna.

Backend core

File	Purpose	Called by	Calls	Edit kab?	Test kaise?
backend/app/main.py	FastAPI app banata hai, middleware order, 7 routers, startup par Qdrant collection ensure	uvicorn	core/* , api/v1/* , vector.ensure_collection	Naya router/middleware	pytest tests/integratio
backend/app/core/config.py	Sabhi settings ek jagah (Settings), .env repo root se padhta hai	Har module	pydantic-settings	Naya config variable	pytest tests/unit + .e
backend/app/core/auth.py	JWT verify (Cognito RS256 ya local dev HS256), tenant + role nikalna	api/deps.py	Cognito JWKS	Auth rules	pytest tests/security/test_plat
backend/app/api/deps.py	CurrentUser , AdminUser , TenantAdminUser , PlatformAdminUser , rate-limited users – har route ki "gate"	Har route	auth.verify_token , ratelimit.enforce	Naya permission type	security tests
backend/app/core/middleware.py	Correlation ID, security headers, trusted hosts, body size limit	main.py	–	Header/limit badalna	integration tests
backend/app/core/ratelimit.py	Redis par per-user rate limit (20/10/5 per min) + per-account auth limit (10/min)	deps.py , api/v1/auth.py	Redis	Limits	security tests
backend/app/services/identity.py	Cognito admin APIs (invite, role/status change, password reset, global sign-out) + \$0 local provider	api/v1/auth.py , admin.py , platform.py	boto3 cognito-idp	Auth ya invite flow	pytest tests/integration/test_c
backend/app/services/onboarding.py	Guard functions – tenant id validation, role assignability, last-admin protection. Koi DB, koi I/O	admin.py , platform.py	–	Hierarchy rules	pytest tests/security/test_onbo

Data layer

File	Purpose	Edit kab?
backend/app/db/models.py	10 tables ka SQLAlchemy model	Naya column/table → saath mein Alembic migration zaroori
backend/app/db/repositories.py	Tenant-filtered data access (documents, jobs, usage, audit, apni company ke users)	Naya query – hamesha <code>tenant_id</code> filter ke saath
backend/app/db/control_plane.py	Ekmatra cross-tenant surface: company registry. Sirf PlatformAdminUser ke peeche, sirf registry data, har mutation audited. repositories.py se alag rakha gaya hai taaki us file ka "sab kuch tenant-filtered hai" contract bina exception bana rahe	Sirf registry feature – yahan koi document/chat query add karna rule violation hai
backend/alembic/versions/0001_initial_schema.py	Poora schema banane wali migration	Kabhi edit mat karo agar kahin apply ho chuki; nayi migration banao
backend/alembic/versions/0002_platform_onboarding.py	PLATFORM_ADMIN enum label, tenants.contact_email, tenants.created_by, users.invited_by, reserved platform tenant row. Sirf additive – purana data chhua nahi jaata	Same

Services (business logic)

File	Kya karti hai
services/rag/pipeline.py	RAG ka dil – 14 stage wala answer_question
services/rag/guardrails.py	Citation validation, output guardrail, confidence
services/rag/cache.py	Redis semantic cache (tenant-namespaced)
services/rag/rerank.py	BM25-style lexical rerank + vector score fusion
services/rag/prompts.py	System prompt (versioned in code), context envelope
services/security/injection.py	Prompt injection scanner (10 pattern families)
services/security/files.py	Upload allow-list, magic bytes, size, safe S3 key
services/ingestion/pipeline.py	Extract → chunk → scan → embed → Qdrant
services/ingestion/extractors.py	PDF, DOCX, CSV, Excel, MD, TXT, image, audio/video
services/ingestion/chunker.py	Paragraph/sentence boundary chunking with overlap
services/vector.py	Qdrant – collection, upsert, mandatory tenant filter search
services/storage.py	S3/MinIO – <code><tenant_id>/<uuid>.<ext></code> keys, presigned URLs
services/ai/provider.py	Bedrock provider + \$0 local provider
services/ai/registry.py	Model roles → model IDs, pricing, vision check
services/ai/transcribe.py	Amazon Transcribe job + polling
services/agents/graph.py, nodes.py	LangGraph agentic workflows

Frontend

File	Purpose	Edit kab?
frontend/src/lib/api.ts	Ekmatra API client – token, silent refresh, correlation ID, errors	Naya endpoint call
frontend/src/components/shell.tsx	Session (email/password sign-in, first-sign-in challenge, password recovery, server-side sign-out), role-aware sidebar nav, page header	Nav item, login screen
frontend/src/components/admin.tsx	Role gates – AdminOnly, TenantAdminOnly, PlatformAdminOnly. Har gate backend ki exactly ek dependency ko mirror karta hai	Naya gated page
frontend/src/app/platform/tenants/page.tsx	Company onboarding wizard + registry (platform operator ke liye)	Onboarding UX
frontend/src/app/admin/users/page.tsx	Apni company ke users – invite, role/department, deactivate, password reset	User management UX
frontend/src/components/ui/index.tsx	Buttons, cards, badges, tables, stat tiles	Design system
frontend/src/components/charts.tsx	SegmentedBar, BarList, RingMeter (no library)	Chart badlav
frontend/src/app/dashboard/page.tsx	Dashboard (real API data)	Dashboard
frontend/next.config.mjs	NEXT_OUTPUT_MODE=export → static; warna standalone	Build mode
amplify.yml	Amplify build commands	Amplify build badalna

Infra & scripts

File	Purpose
infra/terraform/envs/*/main.tf	Chaar Terraform states (Chapter 24)
infra/terraform/modules/service/main.tf	ALB, ECS task (3 containers), CodeDeploy, IAM, alarms
infra/terraform/modules/database/main.tf	RDS PostgreSQL, DB subnet group (write-only password)
infra/terraform/modules/network/main.tf	VPC, subnets, security groups (NO NAT)
infra/terraform/modules/cost-guard/	Budgets, SNS, Lambda (lambda/handler.py), EventBridge
infra/terraform/modules/frontend/main.tf	Amplify app/branch, CloudFront
scripts/_common.sh	Saare scripts ke shared guards (account check, cost guard, reports)
scripts/deploy.sh	Backend deploy pipeline
scripts/deploy-frontend.sh	Frontend deploy + wiring + verification
scripts/bootstrap-platform-admin.sh	Pehla platform operator – seedha Cognito mein. Ekmatra privileged grant jiski koi API nahi hai
scripts/destroy.sh	Ekmatra terraform destroy path

backend/app/workers/ khali hai (sirf __init__.py). Docs "async ingestion workers" bolte hain, lekin code mein ingestion **FastAPI** BackgroundTasks se usi API process ke andar chalta hai (api/v1/documents.py → _run_ingestion_task). Koi alag worker/queue nahi hai.

5. AWS Services – Ye Service Kya Kar Rahi Hai?

Har service ke liye: kyun chahiye · kya store/run karti hai · input/output · kis se judi hai · console mein kahan dekhein · CLI · down ho to kya · cost · interview line.

Region hamesha us-west-2. Resource naam Terraform se hain: ekba-<env>-<component> (yahan env = dev). Account ID ki jagah <ACCOUNT_ID> likha hai.

5.1 Terraform – sab kuch banane wala

Kaam: AWS infrastructure ko code se create/manage karna. Console mein haath se ECS, ALB, S3 banane ke bajaye, Terraform wahi cheezein repeatable tareeke se banata hai. **Kahan:** infra/terraform/ – chaar alag **states** (Chapter 24). **State kahan rehti hai:** S3 bucket ekba-tfstate-<ACCOUNT_ID> + lock ke liye DynamoDB table ekba-tfstate-lock (dono scripts/bootstrap-state.sh banata hai, Terraform ke bahar). **Interview line:** "Infra ko chaar states mein baanta – protected baseline, protected cost guard, ephemeral dev, persistent frontend – taaki destroy kabhi secrets ya ECR images tak pahunch hi na sake."

5.2 S3 – original documents

Kyun	Uploaded original files rakhne ke liye (sirf persistent data store)
Resource	ekba-dev-documents-<random-suffix> (baseline state)
Security	Public access block , SSE AES256, versioning on , TLS-only bucket policy, prevent_destroy
Key format	<tenant_id>/<uuid>.<ext> – server banata hai, filename se kabhi nahi (services/security/files.py → build_storage_key)
Input	Upload API se file bytes
Output	Ingestion ke liye bytes; download ke liye 300 second ka presigned URL
Local mein	S3 ki jagah MinIO container (S3_ENDPOINT_URL=http://localhost:9000)

Console: S3 → Buckets → ekba-dev-documents-... → folders tenant ID ke naam se.

Down ho to: upload aur download fail; baaki search chalta rahega (chunks Qdrant mein hain). Cost: pennies.

AWS par ek important baat (2026-09-12 se badla hai): file S3 mein bachti hai aur uska metadata ab RDS PostgreSQL mein bachta hai – task restart, naya deploy, sab ke baad. Lekin **chunks/vectors (Qdrant)** abhi bhi **ephemeral sidecar** mein hain, isliye task replace hone par vectors chale jaate hain: document list mein dikhega, par uska jawab RAG mein nahi aayega jab tak dobara upload na karo. Seed ke apne documents har startup par khud re-index ho jaate hain. (Re-index job project mein currently implemented nahi hai.)

5.3 Cognito – user identity

Kyun	Login identity + JWT token dena
Resource	User pool ekba-dev-users, app client ekba-dev-web (baseline state)
Username	Email (username_attributes = ["email"])
Custom attributes	custom:tenant_id (1-64 chars), custom:role (1-16 chars) – saara authorization data yahi hai
Password policy	Min 12 chars, upper + lower + number + symbol
MFA	OFF ("demo scope")
App client	Public client (no secret), flows: USER_PASSWORD_AUTH, USER_SRP_AUTH, REFRESH_TOKEN_AUTH
Token validity	Access 1 hour, ID 1 hour, Refresh 7 days
Hosted UI domain	Configure nahi hai – Terraform mein aws_cognito_user_pool_domain nahi hai, aur zaroorat bhi nahi: sign-in apna custom form hai jo /api/v1/auth/login se Cognito ko call karta hai
Custom attributes	custom:tenant_id (max 64), custom:role (max 16 – platform_admin 14 chars mein fit hai)

Console: Cognito → User pools → ekba-dev-users → Users.

Down ho to: naye login nahi ho payenge; backend JWKS se token verify nahi kar payega → 401. Cost: is scale par free tier.

5.4 ECS Fargate – backend chalane wali machine

Kyun	Backend container(s) chalane ke liye bina server manage kiye
Cluster / service	ekba-dev / ekba-dev , launch type FARGATE , deployment controller CODE_DEPLOY
Task size	1 vCPU (1024) / 3 GB (3072) , x86_64
Networking	Public subnet + public IP (NAT Gateway nahi, taaki ECR/Bedrock tak pahunch sake); security group sirf ALB se traffic leta hai
Containers (ek hi task mein)	api (FastAPI), qdrant/qdrant:v1.12.4 , redis:7-alpine – PostgreSQL ab RDS hai (5.17)
api startup command	alembic upgrade head && (python -m seeds.seed echo 'seed failed ...') && exec uvicorn ...
Health checks	api: curl /api/v1/health ; redis: redis-cli ping ; qdrant: koi health check nahi (isliye status UNKNOWN dikhta hai)
Container Insights	disabled (cost bachane ke liye)

Down ho to: poora API down (ALB 503/504). Cost: ~\$0.054/hour (Fargate 1 vCPU/3 GB list price).

5.5 ECR – Docker images ka godown

Repos	ekba-dev-backend (use hota hai), ekba-dev-frontend (bana hai, koi push nahi karta – frontend Amplify se deploy hota hai)
Tags	IMMUTABLE – tag = git commit SHA (7 chars), latest kabhi nahi
Scan	scan_on_push = true ; deploy.sh critical findings par ruk jata hai
Lifecycle	Untagged 1 din mein expire; last 10 images rakho

ECR images manually delete mat karo – CodeDeploy rollback aur redeploy unhi par depend karte hain. ECR repo prevent_destroy aur protected hai.

5.6 ALB – backend ka darwaaza

Resource	ekba-dev-alb (internet-facing, 2 AZ), idle timeout 120 s
Listeners	:80 production, :8080 test (CodeDeploy green ko yahan check karta hai)
Target groups	ekba-dev-blue , ekba-dev-green – har deployment mein roles badalte hain; health check GET /api/v1/health → 200
Security group	Sirf allowed_cidrs (tumhara ek /32 IP) port 80/443 par + CloudFront origin-facing prefix list port 80 par. Port 8080 kisi ke liye open nahi
HTTPS	Nahi – ALB par certificate/HTTPS listener nahi hai; HTTPS CloudFront deta hai

Down ho to: CloudFront 502/504 dega. Cost: ~\$0.023/hour + LCU + 2 public IPv4.

5.7 CloudFront – API ka HTTPS front door

Kyun: Amplify site HTTPS par hai, ALB HTTP-only hai. Browser HTTPS page se HTTP API call block karta hai (mixed content), aur Amplify ka reverse proxy sirf HTTPS targets support karta hai. Isliye CloudFront (default *.cloudfront.net certificate) API ko HTTPS URL deta hai – domain naam ki zaroorat nahi.

Setting	Value (modules/frontend/main.tf)
Cache policy	ekba-dev-api-no-cache – TTL 0 (max 1 s), Authorization cache key mein (GET par forward karne ka yahi tareeka hai; kisi aur user ko response kabhi nahi milega)
Origin request policy	Managed-AllViewerExceptHostHeader – Host ALB ka naam rehta hai, jo TRUSTED_HOSTS (*.elb.amazonaws.com) pass karta hai
Viewer	https-only , PriceClass_100
Origin	ALB DNS, http-only , read timeout 60 s

URL stable rehta hai; sirf origin (ALB naam) deploy-frontend.sh update karta hai. Cost: demo traffic par ~\$0.

5.8 Amplify – frontend hosting

App	ekba-dev-frontend (platform WEB = static hosting), repo Analyst-Parveen/Enterprise-Knowledge-Based-Automation
Branch	aws-deployment , stage PRODUCTION , framework Next.js – SSG , auto-build on push
Build spec	Repo ka amplify.yml (app root frontend , output out/)
Env vars (Terraform)	AMPLIFY_MONOREPO_APP_ROOT=frontend , NEXT_PUBLIC_API_URL=https://<cloudfront> , NEXT_TELEMETRY_DISABLED=1
Custom headers	HSTS, nosniff, X-Frame-Options DENY, Referrer-Policy, CSP (monorepo format applications[].appRoot)
Rule	/<*> → /404.html (404)
GitHub connection	AWS Amplify GitHub App (browser se ek baar authorize) – koi GitHub token Terraform state mein nahi

Console: Amplify → ekba-dev-frontend → aws-deployment → Deployments (har job ke build logs).

Cost: ~\$0.04 per build (~4 build-minutes), idle ~\$0.

5.9 CodeDeploy – blue/green release

App / group	ekba-dev / ekba-dev-dg , config CodeDeployDefault.ECSAllAtOnce
Style	BLUE_GREEN , WITH_TRAFFIC_CONTROL
Rollback window	Purana task set 5 minute (rollback_window_minutes) zinda rehta hai
Auto rollback	DEPLOYMENT_FAILURE aur alarm ekba-dev-5xx par (configure hai, abhi tak exercise nahi hua)
Requirement	Account Paid plan par hona chahiye (Free plan par SubscriptionRequiredException)

5.10 Secrets Manager – secrets

Chaar containers (baseline, recovery_window_in_days = 30 , prevent_destroy): ekba/dev/backend/database-url (JSON: url , password), ekba/dev/backend/qdrant-api-key , ekba/dev/backend/dev-auth-secret , ekba/dev/ai/langsmith-api-key . Values scripts/set-secrets.sh bharta hai – **Terraform state mein kabhi nahi jaate**. ECS execution role task start par inject karta hai. Cost: \$0.40/secret/month ≈ \$1.60/month (baseline ka sabse bada idle kharcha).

5.11 CloudWatch – logs aur alarms

Log group	Retention	Kya hai
/ekba/dev/service	1 din	ECS containers (api , qdrant , redis stream prefixes)
/ekba/dev/audit	30 din	Audit logs (baseline, protected)
/aws/lambda/ekba-dev-cost-guard	14 din	Kill-switch Lambda

Alarms: ekba-dev-5xx (5 se zyada target 5xx/min, 2 minute), ekba-dev-latency (p95 > 10 s). Billing alarm estimated_charges sirf us-east-1 mein banta – is deployment mein **exist nahi karta**; spend budgets cover karte hain.

5.12 Bedrock – AI models

Chat primary us.amazon.nova-lite-v1:0 , fallback us.amazon.nova-micro-v1:0 , vision us.amazon.nova-lite-v1:0 , embeddings amazon.titan-embed-text-v2:0 (1024 dims). Details Chapter 14.

Current limitation: is account ke Bedrock quotas 0 hain (Titan Embeddings V2 requests/tokens per minute, Nova Lite cross-region requests, Nova Lite tokens/day). AWS par har Bedrock call ThrottlingException deta hai – isliye seed fail hota hai aur chat answer nahi deta. Fix: AWS Support case (Service limit increase → Amazon Bedrock, us-west-2).

5.13 Amazon Transcribe

Audio (mp3, wav, m4a, flac, ogg) aur video (mp4, mov, webm) ka transcript banata hai (services/ai/transcribe.py , max wait 15 min). Local mode mein bhi ye AWS call karta hai, isliye seed ka .mp4 locally failed dikhta hai.

5.14 Lambda, SNS, EventBridge, Budgets – cost guard

Poori detail Chapter 30. Short mein: 2 budgets + hourly EventBridge rule → SNS/Lambda ekba-dev-cost-guard → ECS scale to 0 + ALB delete. **Abhi DRY_RUN=true** .

5.15 IAM – kaun kya kar sakta hai

Role	Kaam
ekba-dev-execution	ECS task start: ECR pull, logs, sirf project ke secrets padhna
ekba-dev-task	App runtime: documents bucket objects, Bedrock invoke, Transcribe, aur is project ke apne user pool par invitation flow ke Cognito admin actions (AdminCreateUser , AdminGetUser , AdminUpdateUserAttributes , AdminEnable/DisableUser , AdminResetUserPassword , AdminUserGlobalSignOut). AdminSetUserPassword aur AdminDeleteUser jaan-boojh kar nahi – app kisi ka password choose nahi kar sakta, identity mita nahi sakta
ekba-dev-codedeploy	AWSCodeDeployRoleForECS
ekba-dev-github-deploy	GitHub Actions OIDC deploy role (scoped to this repo)
ekba-dev-cost-guard	Sirf ecs:UpdateService + DeleteLoadBalancer , woh bhi ephemeral tag wale ekba-dev resources par
AmplifySSRLoggingRole-...	Amplify console ne connect ke waqt banaya; static site ke liye use nahi hota; Terraform ignore karta hai

5.16 VPC & security groups

VPC 10.42.0.0/16 – 2 public subnets (ALB + Fargate task, public IP ke saath) aur 2 private subnets (10.42.10.0/24 , 10.42.11.0/24) jinme sirf RDS hai aur koi internet route nahi. Internet gateway hai, **NAT Gateway nahi**. SGs: ekba-dev-alb (ingress: tumhara IP + CloudFront), ekba-dev-tasks (ingress: sirf ALB SG se port 8000/3000), ekba-dev-db (ingress: sirf tasks SG se port 5432, egress kuch nahi).

5.17 RDS – PostgreSQL database (managed)

1. **Kya hai:** AWS ka managed relational database. Pehle Postgres ek sidecar container tha; **2026-09-11 se ye RDS hai.**
2. **Is project mein kya karta hai:** saara relational data rakhta hai – tenants, users, documents, ingestion jobs, conversations, messages, request_usage, feedback, audit events.
3. **Naam:** `ekba-dev-postgres` · `class db.t4g.micro` · PostgreSQL 16 · 20 GB gp3 · single-AZ.
4. **Kahan rehta hai:** VPC ke **private subnets** mein (koi internet route nahi). `publicly_accessible = false`.
5. **Kaun connect kar sakta hai:** sirf backend task. SG `ekba-dev-db` port 5432 sirf `ekba-dev-tasks` SG se allow karta hai – koi CIDR nahi, koi IP nahi. Tumhara laptop seedha connect **nahi** kar sakta (yahi design hai).
6. **Password kahan se aata hai:** Secrets Manager `ekba/dev/backend/database-url` ki `password` key. Terraform use **ephemeral** resource se padh kar **write-only** `password_wo` mein deta hai → password na state mein jaata hai, na plan mein. Container ko wahi `DB_PASSWORD` milta hai aur wo khud `DATABASE_URL` banata hai (`ssl=require`).
7. **Migrations kahan chalti hain:** task ke startup par (`alembic upgrade head`). Schema already sahi ho to ye no-op hai.
8. **Data kab tak bachta hai:** task replace, naya deploy, task crash – sab mein **bachta hai**. `destroy.sh` par snapshot ban jaata hai aur agla `deploy.sh` usi se instance banata hai.
9. **Kharcha:** $\sim \$0.016/\text{hour instance} + \sim \$0.003/\text{hour storage} = \sim \$0.019/\text{hour}$. 24/7 chale to $\sim \$14/\text{month}$ – isiliye ye ephemeral stack ka hissa hai, aur cost guard ise **stop** karta hai (delete kabhi nahi).
10. **Terraform:** `infra/terraform/modules/database`, `envs/dev` se use hota hai. Encrypted at rest, backups 1 din.
11. **Dekhne ka command:**


Dhyan do: Qdrant aur Redis **abhi bhi ephemeral sidecars** hain. Iska matlab: tumhare upload kiye document ki row (Postgres) bach jaati hai, par uske vectors (Qdrant) task replace hone par chale jaate hain – jawab aana band ho jayega jab tak dobara upload na karo. Seed ke documents har startup par khud re-index hote hain.

5.18 Jo services is project mein NAHI hain

Service	Status
Supabase	Project mein currently ye implemented nahi hai
Aurora / RDS Multi-AZ / RDS Proxy	Project mein currently ye implemented nahi hai – ek single-AZ <code>db.t4g.micro</code> hai
ElastiCache	Implemented nahi hai – Redis container
EKS, NAT Gateway, EFS	Rules ke hisaab se kabhi nahi
SQS	Project mein currently ye implemented nahi hai
Route 53 / custom domain / ACM certificate	Project mein currently ye implemented nahi hai
WAF	Project mein currently ye implemented nahi hai

6. Authentication – Login Kaise Kaam Karta Hai

6.1 Flow – ek nazar mein



Pehli baar login karne wale invited user ke liye ek extra step hai: Cognito `NEW_PASSWORD_REQUIRED` challenge bhejta hai, frontend "Set your password" screen dikhati hai, aur `POST /api/v1/auth/new-password` ke baad hi session milta hai. One-time password ek challenge credential hai, session nahi.

Password bhoolne par: `POST /api/v1/auth/forgot-password` (jo hamesha 202 deta hai, chahe email exist kare ya na kare – warna wo account enumeration ban jaata) aur phir code ke saath `POST /api/v1/auth/confirm-password-reset`.

6.2 Basic terms

Cognito – AWS ki user directory. Users, passwords aur tokens sambhalti hai. **JWT** – ek signed string (`header.payload.signature`). Payload mein claims hote hain (`sub`, `email`, `custom:tenant_id`...). Signature ki wajah se koi ise badal nahi sakta. **ID token vs Access token** – Cognito dono deta hai. ID token mein user ke attributes (`custom` attributes bhi) aur `aud` (`client ID`) hota hai. Access token API permission ke liye hota hai aur usme `custom` attributes aam taur par nahi hote.

6.3 Frontend kya karta hai (actual code)

- Sign-in screen (`components/shell.tsx` → `SignIn`) mein **email + password form** hai. Chaar stages: `credentials` → `new-password` (pehli baar) / `forgot` → `reset`.
- Token `sessionStorage` mein jata hai (`lib/api.ts`), refresh token aur expiry ke saath – tab band karo to session gaya. Cookie jaan-boojh kar nahi: ye ek pure SPA hai jo alag API origin se baat karta hai, aur `sessionStorage` khud-ba-khud har cross-site request se bahar rehta hai.
- Har request par header: `Authorization: Bearer <ID token>` + `X-Correlation-ID`.
- **Silent refresh** hai. Token expire hone se 2 minute pehle hi `POST /api/v1/auth/refresh` chal jaata hai, aur 401 aane par ek baar retry hota hai. Ek hi in-flight refresh promise share hota hai, warna 10 parallel requests 10 refresh trigger kar deti.
- Sign-out pehle server par `POST /api/v1/auth/logout` maarta hai (Cognito global sign-out – saare live tokens revoke), phir local session clear karta hai. Sirf local clear karna "logout jaisa lagta hai" hota, logout nahi.
- Refresh bhi fail ho jaye to token delete (`clearToken`) → wapas sign-in screen.
- **"Developer sign-in"** disclosure mein wahi purana token box hai, lekin wo sirf tab render hota hai jab API localhost ho – deployed build mein wo box hota hi nahi.

6.4 Backend kya karta hai (`core/auth.py`)

1. `api/deps.py` → `current_context` Bearer token nikalta hai; nahi hai to 401 "Authentication required."
2. **Local dev mode** (`ENVIRONMENT=dev` aur `DEV_AUTH_ENABLED=true`): token HS256 se `DEV_AUTH_SECRET` ke saath verify (`seeds/dev_token.py` banata hai). Dono conditions na hon to ye path band.
3. **AWS mode**: Cognito JWKS (`https://cognito-idp.us-west-2.amazonaws.com/<POOL_ID>/well-known/jwks.json`) se key lekar RS256 verify: signature, exp, iss, aur aud = `COGNITO_CLIENT_ID`.
4. Claims → `RequestContext`:
 - `user_id = sub`
 - `tenant_id = custom:tenant_id` (ya `tenant_id`) – nahi mila to 401 ("Token is missing a tenant assignment.")

- `role = custom:role (default user); user/admin/platform_admin` ke alawa kuch bhi → 401

5. **Platform role aur platform tenant ka mutual binding:** agar `role == "platform_admin"` hai lekin tenant platform nahi (ya ulta), to 401 + critical security event. Isliye platform identity ka aadha hissa forge karke kuch nahi milta.

6. Ye context poori request mein use hota hai – koi route khud token parse nahi karta.

Errors ka matlab:

Situation	Response	Security event
Token nahi bheja	401 Authentication required.	–
Bekaar/galat token	401 Invalid token.	auth.invalid_token
Expired	401 Token has expired.	auth.token_expired
Galat issuer/audience	401 Invalid token.	auth.bad_issuer / auth.bad_audience
Tenant claim missing	401	auth.missing_tenant
platform_admin normal tenant mein (ya ulta)	401 Token carries an inconsistent tenant assignment.	auth.platform_role_tenant_mismatch (critical)
Admin route par non-admin	403	authz.admin_required
Platform route par tenant role	403	authz.platform_required
Galat email/password	401 Incorrect email or password. – dono cases mein ek hi message	auth.signin_failed
Ek account par 10+ sign-in attempts/min	429 Retry-After ke saath	rate-limit event

6.5 Tenant aur role kahan se aate hain

Sirf verified token se. Request body, query ya header mein `tenant_id` bhejne se kuch nahi hota – ChatRequest mein `tenant_id` field hai hi nahi (schemas.py: "deliberately no tenant_id field").

6.6 ID token vs Access token – kyun ID token

Ye ek interview-favourite detail hai, aur ab isse manually deal karne ki zaroorat nahi – auth endpoints seedha ID token return karte hain.

Wajah: `core/auth.py` `custom:tenant_id` claim maangta hai aur `aud` ko client ID se verify karta hai. Cognito ke standard tokens mein **ye dono ID token** mein hote hain:

	ID token	Access token
<code>custom:tenant_id</code> , <code>custom:role</code>	✓	✗
<code>aud</code> (client ID)	✓	✗ (client_id hota hai, aud nahi)

Matlab access token audience verification par hi fail ho jaata. Agar kabhi "Token is missing a tenant assignment" dikhe, to 90% chance hai ki access token use ho raha hai.

CLI se debug karna ho to:

6.7 Password kabhi app ke paas nahi aata

Ye ek design decision hai, shortcut nahi:

- App password **set, store, read, log ya return** kabhi nahi karta. Kisi bhi API response mein password field nahi hai.
- Naye accounts sirf invitation se bante hain: Cognito khud one-time password generate karke email karta hai, aur invitee pehle sign-in par apna password set karta hai. Share karne layak koi cheez exist hi nahi karti.
- ECS task role ke paas `AdminSetUserPassword` aur `AdminDeleteUser` **jaan-boojh kar nahi** hain – application kisi ka password choose nahi kar sakta, aur identity mita nahi sakta.
- User deactivate karne par live tokens turant revoke hote hain (`AdminUserGlobalSignOut`), expire hone ka intezaar nahi.

6.8 Token expiry

Mode	Validity
Local dev_token	Default 12 ghante (--hours)
Cognito ID token	1 ghanta (app client config)
Refresh token	7 din – frontend ise use karta hai , expiry se 2 min pehle silently renew

7. Company / Tenant / User Creation

7.1 Seedhi baat pehle

Company onboarding ek **product feature** hai, CLI ritual nahi. Teen level hain, aur har level sirf apne se ek neechle wale ko bana sakta hai:

Sideways ya upar ki taraf kuch nahi hota: ek company ka admin na nayi company bana sakta, na platform operator, na doosri company ko chhoo sakta.

7.2 Tenant ID kahan "rehta" hai?

Do jagah, aur dono ka kaam alag hai:

- **Cognito attribute** `custom:tenant_id` – ye **authorization** ka source hai. Token se padha jaata hai, aur sirf yahi count hota hai. Body/query/header mein `tenant_id` bhejne se kuch nahi hota.
- **tenants table (PostgreSQL/RDS)** – ye company ka **record** hai: naam, contact, kisne banaya, kab. Registry yahan hai, permission nahi.

Ye jaan-boojh kar alag hain. Token compromise ho jaye to DB row usse rok nahi sakti, aur DB down ho jaye to authorization tootna nahi chahiye. Isliye tenant ka naam dikhane wala lookup /me mein try/except ke andar hai – DB ka blip ek valid session ko invalid nahi banata.

Tenant ID **permanent** hai: wo S3 prefix, Qdrant retrieval filter aur cache namespace ban jaata hai, isliye baad mein badla nahi ja sakta.

7.3 Example: "ABC Insurance" ko onboard karna

Step 0 – Pehla platform operator (ek hi baar, poore system mein)

Ye ekmatra privileged grant hai jiski koi API nahi hai – jaan-boojh kar. Jo API platform operator bana sakti, wo product ka sabse qeemati target hoti.

Script pehle account, region aur pool ka ownership verify karta hai (naam `ekba-dev-*` aur tag `ProjectCode=ekba`), phir Cognito user banata hai `custom:role=platform_admin` + `custom:tenant_id=platform` ke saath. Password wo set nahi karta – Cognito khud one-time password email karta hai. Dubara chalane par kuch nahi todta.

Step 1 – Company banao (UI se, 30 second)

Platform operator se login karo → **Companies** → **Onboard a company**:

- Naam: `ABC Insurance Private Limited`
- Tenant id: naam se slug apne-aap ban jaata hai (`abc-insurance-private-limited`), aur edit bhi kar sakte ho
- Contact email: optional, sirf record ke liye – ye login **nahi** hai

Reserved ids (`platform`, `admin`, `api`, `root`, `system`...) aur duplicate ids refuse hote hain.

Step 2 – Uska pehla admin invite karo

Usi screen par step 2: email, naam, department. Role **select** karne ko nahi milta – ye screen sirf ek hi kism ka account banati hai (`admin`), taaki galti se platform operator na ban jaye.

Cognito us admin ko one-time password email karta hai. Tum wo password kabhi nahi dekhte.

Step 3 – Admin apne users banata hai

Wo admin login karta hai (pehle sign-in par apna password set karta hai), phir **Users** → **Invite a user**: email, naam, role (`user` ya `admin` – bas), department. Har invite par Cognito phir se one-time password email karta hai.

Step 4 – Isolation apne aap ho jaata hai

ABC ka har upload `abc-insurance.../<uuid>` S3 key, us tenant id ke DB rows, aur us tenant id ke Qdrant points par jaata hai. Northwind ka user – ya Northwind ka **admin** – ise kabhi nahi dekh sakta. Jo platform operator ne ABC banayi, wo bhi ABC ke documents nahi khol sakta.

Role badalna

Company ke apne **Users** page se: role dropdown badlo. Backend Cognito mein mirror karta hai aur `user.role_changed` audit event likhta hai – actor ke saath. CLI se bhi ho sakta hai, lekin phir audit trail nahi banta, isliye UI better hai.

Naya role **agle login** (naye token) se lagta hai.

`custom:tenant_id` **galat** likhna = **galat company ka data access**. UI se banaye gaye users mein ye risk nahi hai (tenant token se aata hai, form se nahi) – lekin CLI se user banate waqt पूरी savdhani. Tenant ID ek baar set ho gaya to use mat badlo, aur ek identity ek hi tenant mein honi chahiye.

7.4 Local mein users (no Cognito)

`seeds/seed.py` ye banata hai:

user_id	Tenant	Role	Email
seed-platform-admin	platform	platform_admin	platform@ekba.example
seed-user-a	seed-tenant-northwind	user	priya@northwind.example
seed-admin-a	seed-tenant-northwind	admin	admin@northwind.example
seed-user-b	seed-tenant-contoso	user	jordan@contoso.example
seed-admin-b	seed-tenant-contoso	admin	admin@contoso.example

Login: inme se koi bhi email + .env ka DEV_AUTH_PASSWORD (default LocalDev!2026). Yahi wahi email/password screen hai jo AWS par chalti hai – local mein LocalIdentityProvider HS256 token mint karta hai, AWS par Cognito RS256 deta hai, aur **application code ek hi rehta hai**.

Token seedha chahiye ho to:

Isse "Developer sign-in" disclosure mein paste karo (sirf localhost par dikhta hai).

Documents sirf northwind mein seed hote hain – contoso admin se login karke tenant isolation demo dikhta hai (documents khali, users list mein sirf contoso ke log, chat "not found").

8. Admin vs User

8.1 Kaun kya kar sakta hai (code ke hisaab se)

Capability	USER	ADMIN	PLATFORM_ADMIN	Code
Apne tenant ke documents dekhna (list/detail/status)	✓	✓	✗ - platform tenant mein documents hi nahi	documents.py (rate-limited)
Document upload	✓	✓	✗	POST /documents - UploadUser (admin-only nahi)
Document download (presigned URL, 5 min)	✓	✓	✗	GET /documents/{id}/download
Document delete	✓ sirf apna (owner)	✓ tenant ka koi bhi	✗	repositories.soft_delete_document
Knowledge chat (RAG)	✓	✓	✗	POST /chat
Agentic workflows	✓	✓	✗	/agents/* (10/min)
Feedback	✓	✓	✗	POST /feedback
AI metrics (/admin/metrics)	✗ 403	✓ apne tenant ka	✗	admin.py
Audit log (/admin/audit)	✗	✓ apne tenant ka	✗	admin.py
Security events (/admin/security)	✗	✓	✓	admin.py
Apni company ka record (/admin/tenant)	✗	✓	✗	admin.py
Apni company ke users - list, invite, role/department, deactivate, password reset	✗	✓	✗ - operator company ki directory se bahar rehta hai	admin.py + TenantAdminUser
Nayi company banana (POST /platform/tenants)	✗	✗ 403	✓	platform.py + PlatformAdminUser
Company ka pehla admin invite karna	✗	✗	✓	platform.py
Company registry dekhna (naam, id, seat counts)	✗	✗	✓	db/control_plane.py
Onboarding trail (/platform/audit)	✗	✗	✓ - sirf tenant.* / user.* / platform.* events	platform.py
platform_admin role kisi ko dena	✗	✗ - schema hi reject karta hai (422), aur server dobara	✗ - iski bhi koi API nahi	onboarding.assert_role_assignable
Doosre tenant ka data	✗	✗ - admin bhi nahi	✗ - operator bhi nahi	tenant-isolation.md

Admin ≠ super-user. Admin ko sirf apne tenant ka operational data milta hai. Frontend button chhupana authorization nahi hai - API khud check karta hai (AdminUser / TenantAdminUser / PlatformAdminUser dependency). Frontend ke teen gates (AdminOnly, TenantAdminOnly, PlatformAdminOnly) exactly inhi teen dependencies ko mirror karte hain, taaki UI aur API kabhi disagree na karein.

Platform operator ≠ super-user bhi. Company onboard karna us company ki chaabi nahi hai. Operator registry dekh sakta hai, content nahi - aur ye rule ka hissa hai, implementation detail nahi.

8.2 Limits (teeno roles par)

Limit	Value	Kahan
API requests	20 / minute / user	Redis ekba:rl:api:<tenant>:<user>
Agent (server) requests	10 / minute	ekba:rl:server:...
Uploads	5 / minute (+ API bucket)	ekba:rl:upload:...
Sign-in / password reset	10 / minute / account	ekba:rl:auth:<sha256 of account>
Upload size	50 MB	S3_UPLOAD_MAX_BYTES
Daily AI spend	\$0.50 / user / din	DAILY_COST_CEILING_USD (request_usage table se)

Auth limit IP par **nahi**, account par hai - CloudFront + ALB ke peeche client IP ya to bahut logon mein shared hoti hai ya client khud bhejta hai, to IP-based limit galat logon ko rokti hai aur asli attacker ke liye bekaar hai. Account identifier Redis tak pahunchne se pehle hash ho jaata hai, isliye kisi key mein email nahi hoti.

8.3 Scenario

9. Frontend – Browser Mein Kya Chal Raha Hai

9.1 Tech stack (frontend/package.json)

Next.js 15.1.3 (App Router) · React 19 · TypeScript (strict) · Tailwind CSS 3.4 · lucide-react icons · koi chart library nahi (apne SVG/CSS charts).

9.2 Do build modes

Mode	Kab	Setting	Output
Static export	AWS Amplify	NEXT_OUTPUT_MODE=export (amplify.yml set karta hai)	frontend/out/ – har page ki HTML file (chat.html , ...)
Standalone server	Docker image (infra/docker/frontend.Dockerfile)	default	.next/standalone/server.js

next.config.mjs export mode mein headers() nahi deta – static site header nahi bhej sakti, isliye security headers **Amplify** set karta hai.

9.3 Pages (routes)

Route	Kaun	Kya dikhata hai
/	–	/dashboard par redirect
/dashboard	sab	Hero, API status, KPI cards, ingestion/content/department charts, recent documents, admin ke liye AI operations
/chat	sab	Knowledge chat – answer, citations, confidence, model, tokens, cost, latency
/agents	sab	Workflows (5 types) chalana
/documents	sab	Upload, list, filter, delete
/departments	sab	Department-wise documents
/usage	sab (data admin ko)	Usage stats
/feedback	sab	Feedback form
/admin/users	admin (tenant admin only)	Apni company ke users – invite, role/department, deactivate/reactivate, password reset
/admin/tenants	admin	"My Company" – company record + seats + documents/chunks
/admin/documents	admin	Read-only document list
/admin/metrics	admin	AI metrics
/admin/security , /admin/audit	admin	Events
/admin/deployments	admin	Blue/green explanation + API health
/platform/tenants	platform_admin	Companies – onboarding wizard (company → pehla admin) + registry, suspend/reactivate
/platform/audit	platform_admin	Onboarding trail – saari companies ke lifecycle events

Nav **role-aware** hai: platform operator ko control plane dikhta hai aur chat/documents **nahi** – kyunki platform tenant mein content hi nahi hai. Ye chhupana security nahi hai; har route server par dobara authorize hota hai.

Saare pages "use client" hain – data browser se API call karke aata hai. Server-side data fetching **nahi hai**, isliye static export sambhav hai.

9.4 API client – src/lib/api.ts

rawRequest (refresh logic ke bina) aur request (refresh ke saath) alag hain – warna refresh call khud refresh trigger karke infinite loop bana deta. FormData body wali request retry nahi hoti, kyunki stream dobara padhi nahi ja sakti.

NEXT_PUBLIC_API_URL :

Environment	Value	Kaun set karta hai
Local	unset → http://localhost:8000	default
AWS Amplify	https://<cloudfront-domain>	Terraform (Amplify env var)

Ye value **build time** par **bundle** mein **bake** hoti hai. Badalne ke baad **rebuild** chahiye. amplify.yml build ko rok deta hai agar value https:// se shuru nahi hoti (mixed content se bachne ke liye).


9.5 Loading / empty / error states

hooks/useAsync.ts har data call ko loading / error / data / reload deta hai. UI mein Skeleton , EmptyState , ErrorState ("Something went wrong" + Try again) components hain.

9.6 Frontend → backend path

10. Backend – FastAPI Andar Se

10.1 Request ka safar (har request)



10.2 Endpoints (sab /api/v1 ke neeche)

Method + path	Auth	Kaam
GET /health	none	Liveness - {"status":"ok","environment": "...", "components": []}
GET /health/ready	none	Postgres + Redis + Qdrant check; koi down → 503
GET /me	user	Token ka principal + company ka naam
POST /auth/login	none (10/min per account)	email + password → ID token, refresh token, ya NEW_PASSWORD_REQUIRED challenge
POST /auth/new-password	none	Pehli baar sign-in - challenge poora karke session
POST /auth/refresh	none	Refresh token → naya ID token
POST /auth/forgot-password	none (10/min)	Reset code email; hamesha 202
POST /auth/confirm-password-reset	none (10/min)	Code + naya password
POST /auth/logout	user	Cognito global sign-out - saare live tokens revoke
POST /documents	user (upload limit)	Upload
GET /documents	user	List (limit ≤ 200)
GET /documents/{id} · /status · /download	user	Detail, ingestion job, presigned URL
DELETE /documents/{id}	owner/admin	Soft delete + Qdrant points + S3 object
POST /chat	user	RAG
POST /feedback	user	Feedback
GET /agents/workflows · POST /agents/run	user (10/min)	Agentic workflows
GET /admin/metrics · /admin/audit · /admin/security	admin	Operations
GET /admin/tenant	admin	Apni company ka record (koi parameter nahi - sirf apni)
GET /admin/users	tenant admin	Apni company ke users
POST /admin/users	tenant admin	User invite (role: user ya admin hi)
PATCH /admin/users/{id}	tenant admin	Role, department, active status
POST /admin/users/{id}/reset-password	tenant admin	Cognito se one-time password email
POST /platform/tenants	platform_admin	Nayi company
GET /platform/tenants	platform_admin	Registry (platform tenant chhod kar)
GET /platform/tenants/{id} · PATCH	platform_admin	Ek company; naam/contact/active status
POST /platform/tenants/{id}/admins	platform_admin	Us company ka pehla admin invite
GET /platform/audit	platform_admin	Onboarding trail (tenant.* / user.* / platform.* hi)

OpenAPI docs: <http://localhost:8000/docs> - sirf dev mein (docs_url dev ke bahar None).

/admin/users ke liye TenantAdminUser chahiye, AdminUser nahi - matlab platform operator ko bhi 403 milta hai. Company onboard karne wala uski user directory se bahar rehta hai.

10.3 Error format

Har error: {"error":{"code":"...", "message":"..."}, "correlation_id":"..."} - stack trace ya SQL kabhi client ko nahi jaata (core/exceptions.py).

11. Database – PostgreSQL

11.1 Technology


PostgreSQL 16, async SQLAlchemy 2.0 + asyncpg, migrations Alembic.

- **Local:** Docker Compose container postgres:16-alpine (volume ke saath, data bachta hai).
- **AWS:** Amazon RDS – ekba-dev-postgres, db.t4g.micro, 20 GB gp3, single-AZ, encrypted, private subnets, SG sirf task SG se. Pehle ye ECS sidecar tha (task restart = khali DB); **ab data bachta hai** – task replace, naya deploy, sab ke baad. destroy.sh snapshot leta hai aur agla deploy.sh usi se restore karta hai.
- **Connection:** container startup par DATABASE_URL khud banta hai – DB_HOST/DB_PORT/DB_NAME/DB_USER (plain env) + DB_PASSWORD (Secrets Manager) + ?ssl=require. Terraform RDS ka password **write-only** argument se deta hai, isliye password kabhi state ya plan mein nahi aata.
- **Migrations:** har task start par alembic upgrade head (schema sahi ho to no-op). Blue aur green **ek hi database** share karte hain, isliye migrations backward-compatible honi chahiye (Chapter 19.3).

11.2 Tables (10)

Table	Kya hai	Tenant column
tenants	Company (id, name, slug, contact_email, created_by)	– (id khud tenant)
users	User (cognito_sub, tenant_id, email, role, department, invited_by)	✓ FK
documents	File metadata, status, chunk_count, S3 source_uri	✓
ingestion_jobs	Processing job (status, progress, error)	✓
conversations	Chat thread	✓
messages	User + assistant messages, citations, model, confidence	✓
request_usage	Har AI call: tokens, cost, latency, cache_hit	✓
user_feedback	Rating (-1/1), reason, comment	✓
prompt_releases	Prompt versions	– (runtime par use nahi hota)
audit_events	Security/audit events (reason codes, payload kabhi nahi)	nullable

11.3 ER diagram (simplified)



Dhyaan do: documents.tenant_id ka tenants par foreign key nahi hai – ye sirf indexed column hai. Chunks database mein nahi, Qdrant mein hain.

11.4 Indexes (tenant filter fast rakhne ke liye)

ix_documents_tenant_status, ix_documents_tenant_department, ix_jobs_tenant_status, ix_conversations_tenant_user, ix_messages_tenant_conv, ix_usage_tenant_created, ix_audit_tenant_created, ix_audit_type, ix_feedback_tenant, uq_users_tenant_email.

11.5 Migrations

Do migrations:

Revision	Kya karta hai
0001_initial_schema	Poora schema – 10 tables, enums, indexes
0002_platform_onboarding	PLATFORM_ADMIN enum label, tenants.contact_email, tenants.created_by, users.invited_by, aur reserved platform tenant row


0002 **poori tarah additive** hai – koi column drop nahi, koi data rewrite nahi, isliye purana tenant/user/document/conversation data waisa hi rehta hai. Do detail dhyaan dene layak:

- ALTER TYPE user_role ADD VALUE IF NOT EXISTS 'PLATFORM_ADMIN' – PostgreSQL enum labels **UPPERCASE** hain (Python values lowercase), aur PG12+ par ye transaction-safe hai jab tak nayi label usi transaction mein use na ho.
- Platform tenant INSERT ... ON CONFLICT (id) DO NOTHING se aata hai, isliye migration dubara chalane par kuch nahi todta. Downgrade us tenant ko sirf tab delete karta hai jab usme koi user na ho.

Naya column: models.py edit → alembic revision --autogenerate -m "..." → generated file review → alembic upgrade head → tests. **Applied migration kabhi edit mat karo.**

12. Document Ingestion Flow

12.1 Diagram



12.2 Har stage

Stage	Input	Kya hota hai	Output	File
Validate	File	Extension allow-list, magic bytes, 50 MB cap, filename sanitize	Safe bytes	security/files.py
Store	Bytes	S3 (local: MinIO) mein <tenant_id>/<uuid>.<ext>	source_uri	storage.py
Record	—	documents (pending) + ingestion_jobs (queued) + audit	IDs	api/v1/documents.py
Extract	Bytes	Type ke hisaab se extractor	Text blocks (page/section ke saath)	ingestion/extractors.py
Chunk	Blocks	Paragraph → sentence boundaries, overlap	Chunks	ingestion/chunker.py
Scan	Chunk text	Hidden instructions dhoondna (retrieval-poisoning defence)	suspicious flag	security/injection.py → scan_content
Embed	Chunks	25-25 ke batch	1024-dim vectors	ai/provider.py
Index	Vectors + payload	Qdrant upsert (tenant_id ke bina refuse)	Points	vector.py
Finish	—	Document ready, job completed, cache invalidate	—	ingestion/pipeline.py

Supported types (files.py): pdf, txt, md, doc, docx, csv, xlsx, xls, png, jpg, jpeg, webp, gif, mp3, wav, m4a, flac, ogg, mp4, mov, webm.

Likely limitation (test nahi kiya): .doc aur .xls allow-list mein hain, lekin extractor python-docx (sirf .docx) aur openpyxl (sirf .xlsx) use karta hai. Purane binary .doc / .xls parse fail hone ki sambhavna hai. requirements.txt ke comment "DOC / DOCX", "XLSX / XLS" isse match nahi karte.

12.3 Status kaise badalta hai

Document: pending → processing → ready (ya failed), delete par deleted. Job: queued → extracting → chunking → embedding → completed (ya failed + error_code). transcribing status enum mein hai.


Failure par: document failed, job failed with error_code (exception ka naam), audit event record hota hai. Background task crash kabhi silently nahi hota - ingestion_task_crashed log hota hai.

12.4 Manually check karna

13. RAG Query Flow – Sabse Important Chapter

13.1 Example

User (northwind) poochta hai: "What is the domestic hotel reimbursement limit?"



13.2 Har step – Hinglish mein

#	Step	Input → Output	File
1	Authentication	Token → RequestContext(user_id, tenant_id, role)	api/deps.py , core/auth.py
–	Daily ceiling	Aaj ka request_usage sum ≥ \$0.50 → "reached today's AI usage limit"	api/v1/chat.py
2	Input validation	Sawaal → cleaned sawaal (max 4000 chars, control chars reject)	security/injection.py → validate_question
3	Injection scan	Sawaal → safe / blocked. Pattern families: instruction_override, system_prompt_extraction, role_manipulation, guardrail_bypass, exfiltration, delimiter_injection, credential_probe, embedded_directive, hidden_instruction_marker, prompt_leak_bait (+ invisible chars, encoded blobs)	injection.py
–	Embed	Sawaal → 1024-dim vector (cache aur search dono ke liye ek hi call)	ai/provider.py
4	Semantic cache	Vector → cached answer agar isi tenant ka koi sawaal cosine ≥ 0.95 ho	rag/cache.py
5–6	Tenant-filtered retrieval	Vector → top 16 chunks, must: tenant_id == ctx.tenant_id, optional document_ids/department filter, suspicious chunks bahar	vector.py → search
7	Relevance threshold	Score < 0.35 (local 0.15) wale hatao; kuch na bache → "not found" answer, model call nahi	rag/pipeline.py
10	Rerank	BM25-style lexical score (weight 0.35) + vector score → top 8	rag/rerank.py
8	Context construction	Chunks → <<<BEGIN_SOURCES>>> ... [S1] ... <<<END_SOURCES>>> data envelope, budget ~8192×4 chars	rag/prompts.py
9	Model call / routing	System prompt + context + sawaal → answer text. Primary fail → fallback (logged, model_used mein dikhta)	ai/provider.py
11	Citation validation	[Sn] markers → sirf retrieved chunks map hote hain; invented hata diye, foreign-tenant citation drop	rag/guardrails.py
12	Output guardrail	System prompt echo, AWS key / private key / JWT, <script>, javascript: ... → block	rag/guardrails.py
–	Confidence	0.6 × avg(top-3 scores) + 0.4 × citation coverage – 0.2 × invented	guardrails.compute_confidence
13	Usage tracking	Tokens, cost, latency → request_usage; clean answer + citations → cache	repositories.record_usage
14	Language	System prompt mein instruction	prompts.py

Phir chat.py conversation + dono messages (user, assistant) DB mein save karta hai.

13.3 Response envelope (frontend ko kya milta hai)

answer, citations[], retrieved_chunks[] (280-char preview), model_used, input_tokens, output_tokens, estimated_cost, latency_ms, cache_hit, tenant_id, confidence, correlation_id, conversation_id

13.4 Hallucination kaise kam hota hai

1. Model ko sirf retrieved context diya jaata hai (system prompt: "grounded in provided context with citations").
2. Relevance threshold – kuch relevant nahi to model call hi nahi, seedha "not found".
3. Retrieved text **data envelope** mein – document ke andar likhi "instructions" follow nahi hoti.
4. Citation validation – banaya hua source number hata diya jaata hai aur confidence gir jaata hai.
5. Low temperature (0.1).

13.5 Local mode vs AWS mode

	Local (AI_PROVIDER=local)	AWS (AI_PROVIDER=bedrock)
Embeddings	Hashed bag-of-words (deterministic, \$0)	Titan Text Embeddings V2
Chat	Extractive stub – context ki lines [Sn] ke saath quote karta hai, model_used = "local-dev-stub"	Nova Lite (fallback Nova Micro)
Threshold	0.15	0.35
Cost	\$0	Per token (Chapter 14)

13.6 Agentic workflows (LangGraph)

POST /api/v1/agents/run – workflows: policy_comparison, summarization, cross_document_analysis, knowledge_extraction, report_generation.
 Graph: plan → retrieve → analyze → synthesize → cite. Bounds: MAX_STEPS = 12, timeout 120 s; timeout par partial result (partial: true).
 Retrieval same tenant filter use karta hai. Input injection scan yahan bhi.

14. AI / LLM / Embeddings

14.1 Models

Role	Model ID	Status
Chat primary	us.amazon.nova-lite-v1:0	Configured (AWS quota 0 → abhi blocked)
Chat fallback	us.amazon.nova-micro-v1:0 (text-only)	Configured, fallback
Vision (images)	us.amazon.nova-lite-v1:0	Configured
Embeddings	amazon.titan-embed-text-v2:0 , 1024 dims, normalized	Configured
Speech	Amazon Transcribe	Configured
Local dev	local-dev-stub (chat + embeddings)	Active locally
OpenAI / Azure / Vertex / Ollama	–	Implemented nahi hai (rules ke hisaab se allowed bhi nahi)
gpt-oss models	–	Plan mein the; us-west-2 mein available nahi the → Nova se replace

Model IDs **configuration** hain (.env / Terraform vars) – badalne ke liye code change nahi. Nova ko **inference-profile ID** (us. prefix) chahiye; bare ID par Bedrock ValidationException deta hai. Titan bina prefix ke.

14.2 Pricing table (code mein, ai/registry.py , per 1M tokens)

Model	Input \$	Output \$
nova-micro	0.035	0.14
nova-lite	0.06	0.24
titan-embed-text-v2	0.02	–

14.3 Flow

Operational limitation (2026-09-11): Bedrock quotas 0. ECS logs mein aata hai bedrock_embed_failed ... ThrottlingException → UpstreamError: The embedding service is unavailable. → seed failed - continuing without demo data . API healthy rehta hai; chat aur ingestion quota milne tak nahi challenge. Quota check:

15. Qdrant – Vector Database

Vector database kya hai? Normal DB exact match dhoondhta hai; vector DB **meaning** se multi-julti cheez dhoondhta hai. **Embedding** = text ka numbers mein roop (yahan 1024 numbers ki list). Milte-julte meaning wale text ke vectors paas-paas hote hain. **Similarity search** = sawaal ke vector ke sabse paas wale chunk vectors dhoondhna (yahan **cosine** distance).

Setting	Value
Image	qdrant/qdrant:v1.12.4
Collection	ekba_chunks (QDRANT_COLLECTION)
Dimension	1024, distance COSINE
Payload index	tenant_id (KEYWORD) – mandatory filter ke liye
Startup	ensure_collection() – collection na ho to banata hai; dimension mismatch par error ("create a new collection and re-index")

Har point ka payload: tenant_id, document_id, chunk_id, document_name, page_number, section, source_uri, owner_id, document_version, created_by, created_at, text, modality, department, suspicious .

Tenant filter: vector.search khud must: tenant_id == ctx.tenant_id banata hai – koi caller filter override nahi kar sakta. `upsert_chunks` bina tenant_id wala chunk likhne se mana karta hai.

Local dekhna: `http://localhost:6333/dashboard` (Qdrant UI) ya `curl -s http://localhost:6333/collections/ekba_chunks` (points_count dekho – seed ke baad 41).

16. Redis – Cache aur Rate Limit

Redis `redis:7-alpine`, **persistence off** (`--save "" --appendonly no`) – restart par sab khali.

Use	Key pattern	TTL
Rate limit	<code>ekba:rl:<bucket>:<tenant_id>:<user_id></code>	60 s window
Semantic cache index	<code>ekba:sc:idx:<tenant_id></code> (last 50 question digests)	3600 s
Semantic cache entry	<code>ekba:sc:e:<tenant_id>:<sha256[:32]></code> (answer, citations, vector, model, confidence)	3600 s

Cache hit = isi tenant ka pehle ka sawaal cosine ≥ 0.95 . **Miss** = poori RAG pipeline. Document upload/delete par us tenant ka cache **invalidate**.

Redis down ho to? Cache ke errors pakde jaate hain (# cache must never break the request path) – chat chalta rahega bina cache ke. Rate limiting Redis par depend karta hai; `/health/ready` Redis ko unhealthy dikhayega.


17. Data Storage Map – Kya Kahan Rehta Hai

Data	Kahan	Kyun	Kaise dekhein
Original files	S3 (local: MinIO) <tenant>/<uuid>.<ext>	Persistent, encrypted	aws s3 ls s3://<BUCKET>/<tenant>/ · MinIO console http://localhost:9001
Document metadata, jobs, conversations, messages, usage, feedback, audit	PostgreSQL – local: container; AWS: RDS (data bachta hai)	Relational	Local: docker exec -it ekba-dev-postgres-1 psql -U ekba -d ekba ; AWS: sirf API/logs se
Chunks + embeddings	Qdrant ekba_chunks	Vector search	http://localhost:6333/dashboard
Cache + rate limit counters	Redis	Speed / abuse control	docker exec -it ekba-dev-redis-1 redis-cli --scan -- pattern 'ekba:*
User identity, tenant, role	Cognito (AWS) / dev token (local)	Auth	Cognito console / list-users
Secrets	Secrets Manager (AWS) / .env (local, git-ignored)	Security	Sirf naam dekho, value nahi
App logs	CloudWatch /ekba/dev/service (AWS) / terminal (local)	Debugging	aws logs tail
Terraform state	S3 ekba-tfstate-<ACCOUNT_ID> + DynamoDB lock	Infra record	terraform state list (kabhi haath se edit nahi)
Deployment reports	docs/reports/*.md (local files)	Audit trail	File kholo (account ID hota hai – commit mat karo)

AWS par PostgreSQL RDS hai (private subnets) aur Qdrant/Redis task ke andar containers hain. Dono tak direct public raasta nahi: RDS ka SG sirf task SG se allow karta hai, aur task ka SG sirf ALB se. AWS par data dekhna ho to API endpoints (admin metrics, documents list) ya CloudWatch logs use karo. ECS Exec configure nahi hai.

18. Local Development – Laptop Par Poora Project

18.1 Picture



Local mein \$0 – koi AWS nahi, local AI stub, local dev tokens.

18.2 Sabse aasaan raasta – ek command

Ye 7 steps karta hai: prerequisites check → .env banana (generated secrets) → backend deps → containers → migrations → seed → frontend deps. Dobara chalana safe hai.

18.3 Haath se, step by step (RUNBOOK Part A se verified)

Step 1 – .env banao

Paanch values set karo: POSTGRES_PASSWORD, DATABASE_URL (same password), DEV_AUTH_SECRET (32+ chars), DEV_AUTH_PASSWORD (jo password se tum local accounts mein login karoge), S3_BUCKET=ekba-dev-documents. Local ke liye ye rehne do: ENVIRONMENT=dev, AI_PROVIDER=local, DEV_AUTH_ENABLED=true, S3_ENDPOINT_URL=http://localhost:9000, RATE_LIMIT_AUTH_ATTEMPTS_PER_MIN=10.

Step 2 – Containers

--env-file .env zaroori hai – compose file infra/docker/ mein hai, bina iske root .env nahi padha jaata aur password mismatch hota hai.

Step 3 – Backend install

Step 4 – Tables → alembic upgrade head Step 5 – Demo data → python -m seeds.seed (output mein chunks_indexed: 41) Step 6 – Backend → uvicorn app.main:app --reload --port 8000 Step 7 – Frontend (naya terminal) → cd frontend && npm install && npm run dev Step 8 – Kholo → http://localhost:3000, aur kisi bhi seeded email + DEV_AUTH_PASSWORD se login karo:

Email	Role
priya@northwind.example	user
admin@northwind.example	company admin
platform@ekba.example	platform operator
admin@contoso.example	doosri company ka admin (isolation demo)

Step 9 (optional) – Token chahiye to → cd backend && python -m seeds.dev_token [--role admin|platform_admin], phir sign-in page par "Developer sign-in" kholo aur paste karo.

18.4 Check, logs, docs

Backend logs uvicorn wale terminal mein aate hain (JSON lines, correlation_id ke saath).


18.5 Stop / restart

Port 8000 atka ho (PowerShell):


Dev server chalte hue `npm run build` mat chalo. Dono `frontend/.next` share karte hain - build dev server ki files overwrite kar deta hai aur pages load hona band ho jaata hai. Pehle dev server roko.

19. AWS Production Flow

19.1 Frontend




19.2 Backend



Backend deploy **tumhare laptop se** hota hai – git push se automatically **nahi** hota.

19.3 Blue/green – kaise kaam karta hai



Har `deploy.sh` ek poora blue/green karta hai (~9 min CodeDeploy, ~11-12 min total), chahe kuch na badla ho. Target groups `ekba-dev-blue` + `ekba-dev-green` har baar roles badalte hain.

Honest note: AppSpec mein lifecycle hooks **nahi** hain – green ka test sirf container health checks + ALB health check (`/api/v1/health`, `liveness`) hai. DB/Qdrant/Redis/RAG smoke test deployment ke dauran nahi hota (rules mein required hai, implemented nahi).

20. "Mujhe Code Change Karna Hai – Ab Kya Karu?"

20.1 Backend change

`deploy.sh` khud karta hai: `cost guard` → `local gate` (dobara) → `Docker build` → `ECR push` (immutable SHA tag; image pehle se ho to reuse) → `scan` → `terraform fmt/validate/plan` → **destroy/delete dikhe to typed confirmation I REVIEWED THIS PLAN** → `apply` → `CodeDeploy` → `verify.sh` → `report`.


`backend/` ya `backend.Dockerfile` mein **uncommitted changes** hon to `deploy.sh` ruk jaata hai – image ka tag commit SHA hai, isliye image mein exactly wahi commit hona chahiye. Pehle commit karo.

Terraform change bhi ho (jaise naya env var): wahi `deploy.sh` – plan mein dikhega (task definition "must be replaced" normal hai, Chapter 24.3).

20.2 Frontend change

Terraform **zaroori nahi** – sirf code change hai to push kaafi hai. Terraform tab chahiye jab Amplify settings/CloudFront badalne hon (`deploy-frontend.sh` wo karta hai).

20.3 Decision tree



21. Branch Strategy – Asli Setup

Branch	Abhi kahan	Kya trigger karta hai
aws-deployment	15c2a00 (latest)	Amplify auto-build (frontend). CI nahi chalta
main	dd545c4 (peeche hai)	CI (ci.yml: push to main/develop, PR to main)
develop	exist nahi karta	CI trigger list mein hai

Backend deploy kisi branch push se nahi – `./scripts/deploy.sh` jo bhi local commit checked-out hai use deploy karta hai. `deploy.yml` (GitHub Actions) manual hai aur abhi tak run nahi hua (Chapter 25).

Tumhara intended workflow: main pe kaam → aws-deployment mein merge → push → Amplify.

Pehle ek baar sync karna padega: main abhi 6 commits peeche hai. Jab tum decide karo:

Uske baad roz ka flow:

`git add .` se bacho – reports (docs/reports/2026*.md, account ID ke saath) untracked padi rehti hain; files naam se add karo.

22. Stop / Start / Destroy – Cheat Sheet

22.1 Stop vs Destroy

	Matlab	Kharcha
STOP	Temporarily band – resource rehta hai	Band cheez ka compute bachta, par ALB jaisi cheezein chalti rehti hain
DESTROY	Resource hi hata do (database ka snapshot pehle ban jaata hai)	envs/dev ka poora ~\$0.11/hr khatam

Project mein AWS "stop" script nahi hai. Supported tareeka destroy.sh hai. ECS ko haath se 0 par scale karna (aws ecs update-service --desired-count 0) ALB ka kharcha nahi rokta (~\$24/month idle) aur Terraform drift banata hai – cost guard (armed hone par) exactly yahi + ALB delete karta hai.

22.2 Commands

Kaam	Command
Local frontend stop / start	Ctrl+C / cd frontend && npm run dev
Local backend stop / start	Ctrl+C / cd backend && uvicorn app.main:app --reload --port 8000
Local containers stop / start	docker compose --env-file .env -f infra/docker/docker-compose.yml stop / up -d ...
AWS kya chal raha hai	./scripts/cost-check.sh
AWS backend deploy / update	./scripts/deploy.sh
AWS frontend deploy / verify	git push origin aws-deployment / ./scripts/deploy-frontend.sh
AWS backend destroy	./scripts/destroy.sh → type DESTROY ekba-dev
Destroy ke baad wapass	./scripts/deploy.sh phir ./scripts/deploy-frontend.sh

22.3 destroy.sh kya karta hai

1. Account, region, workspace, state verify
2. Har resource print karta hai jo hatega
3. Plan mein protected cheez (secrets, ECR, state backend, budget, IAM user...) → **abort**
4. Typed confirmation DESTROY ekba-dev (sirf y nahi)
5. Sirf envs/dev state destroy
6. Secrets bache hain confirm + audit report

Bachta hai: baseline (ECR images, S3 documents, Cognito, secrets, budget, audit logs, OIDC role), cost guard, frontend (Amplify + CloudFront), state backend. **Hatta hai:** VPC, subnets, SGs, ALB, target groups, ECS cluster/service/task defs, CodeDeploy app, task IAM roles, service log group, alarms.

22.4 Destroy ke baad redeploy

Frontend aur API URL **same** rehte hain.

22.5 Rollback – galat version live ho gaya to kya karu?

Kya	Kaise	Detail
Frontend – pichla build wapas	./scripts/rollback-frontend.sh	Live build se pehle wala last successful build dobara release karta hai (Amplify job)
Frontend – specific commit	./scripts/rollback-frontend.sh --to <COMMIT_SHA>	Us commit ka build
Frontend – rollback permanent karna	git revert <BAD_SHA> → git push origin aws-deployment	Agla push branch head ko hi deploy karta hai, isliye bina revert ke rollback agle push par hat jaata hai
Backend – automatic	CodeDeploy khud	Deployment fail (DEPLOYMENT_FAILURE) ya alarm ekba-dev-5xx → traffic purane (blue) task set par. Purana set 5 minute zinda rehta hai. Configured hai, abhi tak exercise nahi hua
Backend – deployment chal raha ho	aws deploy stop-deployment --deployment-id <DEPLOYMENT_ID> --auto-rollback-enabled	Sirf in-flight deployment ke liye
Backend – deploy ho chuka, galat nikla	git revert <BAD_SHA> → commit → ./scripts/deploy.sh	"Fix forward": naya image, naya blue/green deploy. Recommended tareeka

./scripts/rollback.sh **traffic wapas NAHI shift karta**. Uska CodeDeploy step code mein "not yet implemented (Phase 5)" hai – ye sirf verify.sh chalata hai aur report likhta hai. Isliye backend rollback = CodeDeploy auto-rollback ya git revert + deploy.sh .
Purana commit git checkout karke deploy.sh chalana possible hai (same SHA ka image ECR mein ho to reuse hota hai), lekin us commit ka **Terraform bhi apply hoga** – purana Terraform naye infra se mismatch kar sakta hai. Isse bacho.

AWS frontend "stop" karna: Project mein currently ye implemented nahi hai – Amplify ka koi stop button/script nahi hai aur destroy.sh frontend ko touch nahi karta. Zaroorat bhi nahi: idle Amplify + CloudFront ka kharcha ~\$0 hai. Backend destroy ho to site khulti rahegi par "Could not reach the API" dikhayegi. envs/frontend par seedha terraform destroy project rules ke khilaaf hai (destroy sirf destroy.sh se).

23. Docker

Image = packaged app (code + dependencies). **Container** = image ka chalta hua instance. **Port** = bahar se andar ka raasta (-p 8000:8000). **Volume** = container ke bahar data (restart par bachta hai). **Network** = containers aapas mein baat karein. **Compose** = kai containers ek file se.

File	Kya
infra/docker/docker-compose.yml	Local: postgres, qdrant, redis, minio, minio-init, + optional api aur web
infra/docker/backend.Dockerfile	python:3.12-slim, non-root user ekba (uid 10001), port 8000, HEALTHCHECK, uvicorn
infra/docker/frontend.Dockerfile	node:22-alpine, Next.js standalone (Amplify ise use nahi karta)

AWS par Docker Compose **nahi** – wahi chaar images ECS task definition mein containers ban kar chalti hain.

24. Terraform – Beginner Guide

24.1 Basic commands

Command	Matlab
terraform init -backend-config=backend.hcl	Providers download + remote state se jodna
terraform fmt -recursive	Code format
terraform validate	Syntax/logic check (AWS ko touch nahi karta)
terraform plan -out=tfplan	"Kya badlega" – padho!
terraform apply tfplan	Sirf wahi reviewed plan lagao
terraform destroy	Is project mein sirf scripts/destroy.sh ke through

State = Terraform ki diary ki usne kya banaya. **Backend** = state kahan rehti (S3 + DynamoDB lock). **tfvars** = real values (git-ignored). **Module** = reusable block (modules/). **Environment** = envs/<name> – har ek ki alag state.

24.2 Chaar states

State	Key	Kaun chalata hai	Destroy?
envs/baseline	baseline/terraform.tfstate	Haath se, ek baar	Kabhi nahi
envs/cost-guard	cost-guard/terraform.tfstate	Haath se	Kabhi nahi
envs/dev	dev/terraform.tfstate	Sirf deploy.sh (image + snapshot var deta hai)	Sirf destroy.sh (pehle DB snapshot)
envs/frontend	frontend/terraform.tfstate	Sirf deploy-frontend.sh (ALB + app ID var deta hai)	Kabhi nahi (script nahi)

Har folder mein git-ignored backend.hcl (account ID) aur terraform.tfvars (real values); commit hota hai sirf .example.

24.3 "Plan mein destroy/replace dikha – ab kya karu?"

Kabhi blindly apply mat karo. Checklist:

1. **Kaunsa resource?** terraform show tfplan | grep -E "destroyed|replaced"
2. **Kyun?** Us resource ke andar # forces replacement wali line dhoondho – wahi attribute wajah hai.
3. **Kya ye expected hai?** Neeche table dekho.
4. **Data jayega?** S3, secrets, Cognito, ECR – inka replace kabhi accept mat karo.
5. **Unsure ho to ruko.** Plan file delete karo, kuch apply mat karo.

Plan mein	Normal hai?	Kyun
aws_ecs_task_definition.app must be replaced	✅ Haan	Task definition revisions immutable hain; env var/image change = nayi revision
aws_security_group.alb updated in-place	✅	Ingress rule change
aws_amplify_app ... must be replaced	❌ NAHI	Real example is project se: console ne iam_service_role_arn lagaya tha; Terraform use hataana chahta tha → provider app replace karta (GitHub connection khatam). Fix: lifecycle { ignore_changes = [iam_service_role_arn] }
Kuch bhi destroy baseline mein	❌	Baseline protected hai

24.4 State lock atka ho

Error acquiring the state lock → pehle check karo koi aur terraform/deploy chal to nahi raha. Sach mein atka ho tab hi terraform force-unlock <LOCK_ID> – ye human decision hai, rule ke hisaab se record karo.

25. CI/CD – GitHub Actions

25.1 ci.yml (automatic)

Trigger: push to main/develop, pull request to main. (aws-deployment par nahi.)

Job	Steps	Blocking?
backend	ruff lint, mypy, pytest unit/integration/security/evaluation	mypy non-blocking (warning)
frontend	npm ci, typecheck, lint, build, static-export build	Blocking
security	gitleaks secret scan, pip audit, npm audit, checkov (Terraform)	npm audit warning; checkov soft_fail
terraform	fmt -check, validate baseline/dev/cost-guard/frontend	Blocking
docker	backend image build, Trivy scan, non-root check	Blocking

25.2 deploy.yml (manual, workflow_dispatch)

Type DEPLOY → OIDC role → build/push → plan → reject destroys → apply → CodeDeploy → health → verify → rollback on failure.

Abhi tak run nahi hua, aur 3 known problems hain: (1) iska cost check credits ke baad ka spend padhta hai aur deploy role ke पास ce:GetCostAndUsage nahi – kabhi block nahi karega; (2) GitHub runner ka IP DEMO_ALLOWED_CIDR mein nahi – health check ALB tak nahi pahunchega (CloudFront URL use karna fix hoga); (3) allowed_cidrs do jagah se set hota hai. Backend ke liye ./scripts/deploy.sh use karo.


Required GitHub settings (deploy.yml ke liye): repo variables AWS_REGION, EXPECTED_AWS_ACCOUNT_ID, DEMO_ALLOWED_CIDR, S3_BUCKET, COGNITO_CLIENT_ID; secret AWS_DEPLOY_ROLE_ARN. Ye set hain ya nahi – repository se confirm nahi ho paya.

26. Testing

Kya	Command	Kahan se
Backend lint	<code>ruff check app tests seeds · ruff format --check app tests seeds</code>	backend/
Unit	<code>pytest tests/unit -q</code>	mocks, no services
Integration	<code>pytest tests/integration -q</code>	API smoke, RAG pipeline
Security	<code>pytest tests/security -q</code>	tenant isolation, guardrails, platform security, agent isolation, role hierarchy
Evaluation	<code>pytest tests/evaluation -q</code>	RAG quality harness
Onboarding journey	<code>pytest tests/e2e -q</code>	Do companies onboard karke isolation prove karta hai (PostgreSQL chahiye)
Sab ek saath (deploy gate)	<code>ENVIRONMENT=dev AI_PROVIDER=local DEV_AUTH_ENABLED=true pytest tests/unit tests/integration tests/security tests/evaluation -q → 279 tests</code>	deploy.sh
Frontend	<code>npm run lint · npm run typecheck · npm run build</code>	frontend/
Amplify jaisa build	<code>NEXT_OUTPUT_MODE=export NEXT_PUBLIC_API_URL=https://x.invalid npm run build</code>	frontend/
E2E (Playwright, 29 journeys)	<code>./scripts/test-e2e.sh (tokens khud banata hai) ya npx playwright test</code>	local stack chahiye
Cost-guard Lambda	<code>backend/.venv/Scripts/python.exe -m pytest infra/terraform/modules/cost-guard/lambda/tests -q (27)</code>	repo root
Terraform	<code>terraform fmt -check -recursive infra/terraform · terraform -chdir=.. validate</code>	

Rule: test ko kabhi weak/delete mat karo build pass karane ke liye – code fix karo.

27. Monitoring & Logs – "Problem Aaya To Sabse Pehle Kahan Dekhu?"




Correlation ID: har request ka `X-Correlation-ID` frontend banata hai; backend har log line, error response aur audit event mein wahi ID rakhta hai. Error ke `correlation_id` se logs mein search karo.

Aur jagah: admin UI (AI Metrics, Security, Audit pages), `audit_events` table, `./scripts/cost-check.sh`, CloudWatch alarms `ekba-dev-5xx` / `ekba-dev-latency`.

Monitoring ki limitations: custom CloudWatch metrics (tokens, cost, cache hits as metrics) aur LangSmith tracing – **implemented nahi hain**; ye data `request_usage` table aur admin metrics endpoint mein hai.

28. Troubleshooting Playbook

28.1 "Kuch toot gaya" – decision tree



28.2 Problem → Check → Command → Wajah → Fix

Frontend / Amplify

Problem	Check / Command	Possible cause	Fix
Frontend nahi khul raha	<code>aws amplify list-jobs --app-id <APP_ID> --branch-name aws-deployment --query 'jobSummaries[0].[jobId,status]'</code>	Latest build failed	Console → job → BUILD logs
Build fail: NEXT_PUBLIC_API_URL must be an https:// URL, got ''	Build log	App par env var set nahi (console se naya app banane par pehli build)	<code>./scripts/deploy-frontend.sh</code> (Terraform env var set karta hai)
Build fail: Monorepo spec provided without "applications" key	Build log (build "completed successfully" ke baad)	Custom headers galat format mein	<code>modules/frontend/main.tf</code> – headers <code>applications[].appRoot</code> ke neeche (fix commit 6a3e2ef)
Page khulta hai par "Could not reach the API"	Browser DevTools → Network	API down / galat <code>NEXT_PUBLIC_API_URL</code> / CORS	Neeche API aur CORS rows
Unknown URL par 200 aata hai, 404 nahi	<code>curl -sIL <SITE>/nope</code>	Amplify ka behaviour: 301 → 302 → /404.html → 200	Expected – visitor ko not-found page dikhta hai

API path (CloudFront → ALB → ECS)

Problem	Check / Command	Possible cause	Fix
CloudFront 504	<code>curl -s -o /dev/null -w '%{http_code}' https://<CF>/api/v1/health</code>	ALB SG CloudFront ko allow nahi karta	<code>deploy-frontend.sh</code> (wiring) / <code>deploy.sh</code>
CloudFront 502/504, ALB gayab	<code>aws elbv2 describe-load-balancers --names ekba-dev-alb</code>	<code>destroy.sh</code> ya armed cost guard ne ALB hataya	<code>./scripts/destroy.sh</code> (cleanup) → <code>deploy.sh</code> → <code>deploy-frontend.sh</code>
ALB targets unhealthy	<code>aws elbv2 describe-target-health --target-group-arn <TG_ARN></code>	Container crash / health path	ECS events + logs; health path <code>/api/v1/health</code> TrustedHost se exempt hai
Target unused mid-deploy	same	Us group par abhi listener nahi	Normal jab tak <code>AllowTestTraffic</code>
ECS task stopped / baar-baar restart	<code>aws ecs describe-tasks --cluster ekba-dev --tasks <TASK_ARN> --query 'tasks[0].[stoppedReason,containers[]].[name,exitCode,reason]'</code>	Secret missing, image pull fail, essential container (redis/qdrant) down, ya RDS reachable nahi	CloudWatch <code>/ekba/dev/service</code> stream prefix dekho
<code>verify.sh</code> : FAIL health endpoint, <code>curl</code> exit 28	<code>curl -s https://checkip.amazonaws.com</code> vs <code>grep allowed_cidrs infra/terraform/envs/dev/terraform.tfvars</code>	Tumhara IP badal gaya (ISP)	<code>allowed_cidrs</code> update → <code>deploy.sh</code> (1 in-place SG change)
CORS error browser mein	<code>curl -s -D - -o /dev/null -X OPTIONS https://<CF>/api/v1/me -H "Origin: <SITE>" -H "Access-Control-Request-Method: GET"</code>	Live task definition mein origin nahi	<code>deploy-frontend.sh</code> (<code>frontend.auto.tfvars</code>) → <code>deploy.sh</code>
400 Bad Request har call par	Local: <code>.env</code>	Host <code>TRUSTED_HOSTS</code> mein nahi	Host add karo

Deploy

Problem	Check / Command	Possible cause	Fix
SubscriptionRequiredException (CodeDeploy)	<code>aws freetier get-account-plan-state --region us-east-1</code>	Account Free plan par	Paid plan (IAM wajah nahi)
<code>deploy.sh</code> refuse: uncommitted image inputs	<code>git status -- backend infra/docker/backend.Dockerfile</code>	Image tag commit se match nahi karega	Commit karo
<code>deploy.sh</code> refuse: gross usage ≥ \$18	<code>./scripts/cost-check.sh</code>	Cost guard threshold	Destroy + review spend
Docker build fail	<code>docker version</code>	Docker Desktop band	Start karo
ECR push fail	<code>aws ecr get-login-password ...</code>	Login / permission	<code>deploy.sh</code> khud login karta hai; credentials check
Image scan critical	<code>aws ecr describe-image-scan-findings --repository-name ekba-dev-backend --image-id imageTag=<SHA></code>	Base image CVE	Base image update, rebuild
CodeDeploy failed	<code>aws deploy get-deployment-target --deployment-id <ID> --target-id ekba-dev:ekba-dev --query 'deploymentTarget.ecsTarget.lifecycleEvents[].[lifecycleEventName,status]' --output text</code>	Green unhealthy	Logs; auto-rollback configured
DeploymentTargetDoesNotExistException	-	ECS target ID format galat	<code>--target-id ekba-dev:ekba-dev</code>
<code>deploy.sh</code> : rollback also failed - ESCALATE TO A HUMAN	<code>scripts/rollback.sh</code> padho	<code>rollback.sh</code> ka CodeDeploy traffic shift implemented nahi hai ("not yet implemented (Phase 5)") - ye sirf <code>verify.sh</code> dobara chalata hai. Asli rollback CodeDeploy ka auto-rollback karta hai	Asli failure (aksar IP) diagnose karo

Auth

Problem	Check	Possible cause	Fix
"Incorrect email or password."	Backend log <code>auth.signin_failed</code>	Account exist nahi karta, ya password galat – message jaan-boojh kar dono ke liye same hai (account enumeration se bachne ke liye)	Local: <code>python -m seeds.seed</code> chalao aur <code>DEV_AUTH_PASSWORD</code> use karo. AWS: Cognito mein user status check karo
"Too many attempts. Try again shortly."	Retry-After header	Ek account par 10 sign-in attempts/min	Ek minute ruko. Limit IP par nahi, account par hai
Pehli baar login par "Set your password" screen	–	Ye expected hai – invitation ka one-time password ek challenge credential hai, session nahi	Naya password set karo (12+ chars, mixed case, number, symbol)
"Token carries an inconsistent tenant assignment."	Backend log <code>auth.platform_role_tenant_mismatch</code> (critical)	Token mein <code>platform_admin</code> + normal tenant, ya normal role + platform tenant	Expected behaviour. <code>--role platform_admin</code> bina <code>--tenant</code> ke use karo
"Token is missing a tenant assignment."	Token decode (<code>jwt.io</code> jaisa tool mat use karo real token ke liye; locally <code>python -c se</code>)	Access token use hua / user par <code>custom:tenant_id</code> nahi	ID token use karo (Chapter 6.6); attribute set karo
"Invalid token."	Backend log <code>auth.invalid_token</code> / <code>auth.bad_audience</code> / <code>auth.bad_issuer</code>	Galat pool/client, dev token AWS par	<code>COGNITO_USER_POOL_ID</code> / client ID match
"Token has expired."	–	1 ghanta (Cognito) / 12 ghante (dev). Real session khud renew hoti hai	Refresh token gaya to dobara login
"Authentication is not configured."	–	<code>COGNITO_USER_POOL_ID</code> khali aur dev auth off	Env set karo
403 admin page	–	Role <code>user</code>	<code>custom:role=admin</code> , naya login
Companies pages nahi dikh rahe	<code>/api/v1/me</code> ka role	Tum admin ho, <code>platform_admin</code> nahi	Platform operator se login karo
Platform operator ko <code>/admin/users</code> par 403	–	Expected – operator company ki user directory se bahar rehta hai	Company ke apne admin se login karo
Invite bheja par email nahi aaya	Cognito → Users mein status <code>FORCE_CHANGE_PASSWORD</code> ?	Cognito default email sandbox limits; SES configure nahi	Cognito console se resend; production ke liye SES

Data / RAG

Problem	Check / Command	Possible cause	Fix
DB connection fail (local)	<code>curl localhost:8000/api/v1/health/ready</code>	<code>DATABASE_URL==...</code> double =, password mismatch	<code>.env</code> fix; mismatch par <code>docker compose ... down -v</code> (data jayega)
Qdrant fail	<code>curl localhost:6333/collections/ekba_chunks</code>	Container down / dimension mismatch	Up karo; dimension badli ho to nayi collection + re-index
Redis fail	<code>docker exec -it ekba-dev-redis-1 redis-cli ping</code>	Container down	Cache bina chalega; rate limit ke liye Redis chahiye
Upload 415	Response	Extension/magic bytes mismatch	Sahi file type
Upload 413	Response	> 50 MB	Chhoti file
Upload 429	Response Retry-After	5 uploads/min	Ruko
Document processing mein atka	<code>/documents/<id>/status</code> , logs <code>ingestion_task_crashed</code>	Bedrock throttling; Transcribe (max 15 min); container restart ne background task maar diya (ingestion same process mein chalti hai)	Logs dekho; dobara upload
RAG "I could not find anything..."	Doc ready? <code>points_count</code> ?	Nothing above threshold (0.35/0.15), department filter, doc processing, Bedrock quota	Doc status + Qdrant check
<code>bedrock_embed_failed ... ThrottlingException</code>	Service Quotas	Quota 0	AWS Support case
Galat tenant ka data dikha	Audit <code>tenant.cross_tenant_*</code> events	Critical incident	Turant roko; token claims check; <code>pytest tests/security -q</code>

Terraform / cost

Problem	Fix
Error acquiring the state lock	Koi aur run chal raha? Nahi to hi <code>terraform force-unlock <LOCK_ID></code> (decision record karo)
Plan mein unexpected destroy/replace	Chapter 24.3 checklist – apply mat karo
Cost badh raha hai	<code>./scripts/cost-check.sh</code> ; bhoola hua stack → <code>./scripts/destroy.sh</code>

Windows shell traps (asli mein hue):

- Git Bash `/ekba/dev/service` ko Windows path bana deta hai → `MSYS_NO_PATHCONV=1` aws logs ...
- PowerShell 5.1 `-backend-config=backend.hcl` ko dot par tod deta hai → quote karo ya Git Bash use karo
- AWS CLI `--max-items` ke saath `--output text` extra None line deta hai (Windows par `\r` bhi) → scripts mein `--max-items` mat lagao, `tr -d '\r'`

29. Security

29.1 Controls – kya kahan

Control	Implementation	File
Authentication	Cognito RS256 JWT (signature, issuer, audience, expiry); dev HS256 sirf dev mein. Sign-in email+password, invitation-only accounts, generic failure message	core/auth.py , api/v1/auth.py , services/identity.py
Authorization	Roles user / admin / platform_admin ; AdminUser / TenantAdminUser / PlatformAdminUser dependencies	api/deps.py
Privilege escalation	platform_admin kisi bhi request schema mein valid value nahi (422), aur assert_role_assignable server par dobara refuse karta hai. Platform role aur platform tenant token verify par ek doosre se bandhe hain	schemas.py , services/onboarding.py , core/auth.py
Cross-tenant admin view	Ekmatra jagah: company registry. Sirf PlatformAdminUser , sirf registry data, har mutation audited	db/control_plane.py
Credential handling	App password set/store/log/return kabhi nahi karta; Cognito one-time password email karta hai; deactivate par live tokens revoke	services/identity.py
Lockout protection	Admin apna role/status nahi badal sakta; company ka last active admin nahi hataya ja sakta	services/onboarding.py
Tenant isolation	Token se tenant; har DB query, Qdrant search (mandatory filter), S3 prefix check, cache key mein tenant	repositories.py , vector.py , storage.py , cache.py
Cross-tenant attempt	404/403 response + critical audit event tenant.cross_tenant_document_access	repositories.require_document
Prompt injection	10 pattern families + invisible chars + encoded blobs; block par audit	security/injection.py
Retrieval poisoning	Ingestion par har chunk scan; suspicious chunks search se bahar	ingestion/pipeline.py , vector.py
Output guardrail	System prompt leak, AWS key, private key, JWT, script/iframe/js URI block	rag/guardrails.py
Citation validation	Invented/foreign-tenant citations hatao	rag/guardrails.py
File upload	Allow-list, magic bytes, 50 MB, server-generated key, sanitized name	security/files.py
Rate limits	20 / 10 / 5 per minute per user, aur 10 sign-in attempts/min per account (hashed key) → 429	core/ratelimit.py
Headers (API)	Security headers middleware	core/middleware.py
Headers (site)	HSTS, nosniff, DENY, no-referrer, CSP	Amplify (modules/frontend)
CORS	Explicit allow-list (never *)	main.py + CORS_ALLOWED_ORIGINS
Trusted hosts	TRUSTED_HOSTS (health path exempt for ALB)	core/middleware.py
Network	Tasks sirf ALB se; ALB sirf tumhara IP + CloudFront	modules/network
S3	Public block, SSE, versioning, TLS-only	baseline
Secrets	Secrets Manager, runtime injection, SecretStr , never in state/logs	baseline, config.py , logging.py (redaction)
IAM	Role per component, scoped ARNs	modules/service , baseline
Container	Non-root user (uid 10001)	backend.Dockerfile
Audit	audit_events (reason codes, payload kabhi nahi)	repositories.record_audit

29.2 Honest gaps

Gap	Detail
ALB par HTTPS nahi	Browser traffic CloudFront par HTTPS hai; CloudFront→ALB HTTP
API CloudFront se public	Ab kahin se bhi CloudFront ke through reachable – /health ke alawa sab ke liye valid token chahiye
MFA off	Cognito mfa_configuration = "OFF"
CSP 'unsafe-inline'	Next.js inline bootstrap scripts ke liye
Next.js 15.1.3 advisory	Amplify build CVE-2025-66478 warning dikhata hai (server rendering se related; site static hai) – upgrade alag decision
npm audit	7 vulnerabilities (Amplify log)
WAF	Implemented nahi
Custom domain / ACM	Implemented nahi

29.3 GitHub mein kabhi commit nahi karna (.gitignore se)

.env , .env.* (sirf .env.example allowed) · *.pem *.key *.p12 *.pfx *.crt · credentials , aws-credentials , secrets.json , .secrets/ · *.tfstate* , tfplan , *.tfplan , tfdestroyplan · *.tfvars (sirf *.tfvars.example) · .terraform/ · backend.hcl (account ID) · envs/dev/frontend.auto.tfvars (*.tfvars pattern) · aur convention se: docs/reports/2026*.md (account ID + live URLs hote hain).


`.env.example` mein sirf placeholders (`<replace-me>`). CI mein gitleaks secret scan chalta hai.

30. Cost Guard – Kharcha Kaise Control Hota Hai

30.1 Pehle ek sach

AWS Budget koi guaranteed hard spending cap nahi hai. Billing data late aata hai (Budgets din mein kuch baar refresh, data ~24 ghante tak late). Isliye shutdown \$18 par set hai, \$20 ceiling se \$2 neeche – asal mein shutdown ~\$18-20 ke beech hota hai.

30.2 Architecture



30.3 Thresholds

Trigger	Threshold	Action
Gross usage (credits + refunds excluded), cumulative	> \$10, > \$15	Email
Same	Forecast > \$18	Email
Same	Actual > \$18	Email + shutdown
Net spend (credits ke baad = card charge), monthly	> \$0.01	Email + shutdown
Hourly check	Environment 8 ghante se purana (ALB/service creation se, redeploy reset nahi karta)	Shutdown (dry-run mein sirf log)

deploy.sh bhi \$18 gross par deploy refuse karta hai (COST_GUARD_SHUTDOWN_USD, COST_GUARD_START=2026-09-01).

30.4 Shutdown kya karta hai / kya chhodta hai

Karta: ECS service ekba-dev → 0 tasks (delete nahi) · ALB ekba-dev-alb + listeners **delete** (ALB pause nahi hota; idle ALB ~\$24/month). **Chhodta:** VPC, subnets, SGs, target groups, cluster, task defs, IAM, logs, alarms (\$0 idle – destroy.sh hatata hai) · poora baseline · frontend · state backend. **Kabhi terraform destroy nahi chalata.**

Safety: har resource ke liye teen checks – naam ekba-dev/ekba-dev-*, tags ProjectCode=ekba + Environment=dev + Lifecycle=ephemeral , aur ARN s3://ekba-tfstate-<ACCOUNT_ID>/dev/terraform.tfstate mein. IAM role bhi sirf yahi do actions allow karta hai (policy simulator se verified).

30.5 DRY_RUN

- **Abhi:** true (live Lambda, committed example, local tfvars – teeno).
- **Fail-safe:** false ke alawa koi bhi value = dry run.
- Dry-run mein hourly check sirf log karta hai (email spam nahi); manual/budget trigger email bhejte hain.

Arm karna (jab tum decide karo):

Git push se ye apply nahi hota – terraform.tfvars git-ignored hai aur koi workflow cost-guard apply nahi karta. Arm karte hi agli hourly run 8 ghante se purana stack band kar degi. Demo chal raha ho to pehle destroy.sh + deploy.sh (fresh 8 ghante).

30.6 Manually verify

Dry-run result (2026-09-12): WOULD scale ECS service ekba-dev/ekba-dev from 1 to 0 tasks · WOULD delete load balancer ekba-dev-alb (and its listeners) · WOULD stop RDS instance ekba-dev-postgres (status available; data and storage kept) · refused 0 · errors 0.

Shutdown ke baad: ./scripts/destroy.sh → ./scripts/deploy.sh → ./scripts/deploy-frontend.sh (sirf deploy.sh service ko 0 par chhod dega kyunki Terraform desired_count ignore karta hai).

30.7 Kharcha – numbers

Cheez	Cost
Backend stack chalte hue	~\$0.11/hour (Fargate ~\$0.054, ALB ~\$0.023 + LCU, 3 public IPv4 ~\$0.015, RDS ~\$0.019) → 4 ghante ~\$0.44, mahina bhool gaye to ~\$80
RDS akela, 24/7	~\$14/month (\$11.68 instance + \$2.30 storage) – isiliye session ke baad destroy.sh
Database snapshots (destroy ke beech)	~\$0.095 per GB per month – yahan kuch cents
Baseline idle	~\$1.70/month (4 secrets × \$0.40)
Cost guard	~\$0 (free tiers)
Amplify	~\$0.04/build, idle ~\$0
CloudFront	demo traffic par ~\$0
Cost Explorer API	\$0.01 per call (deploy.sh , cost-check.sh)

Baseline budget ekba-dev-monthly (\$20, 50/80/100%) credits **include** karta hai → credits khatam hone tak \$0 dikhata hai. Asli protection upar wale gross/net budgets hain.

31. Manual Data Inspection (read-only)

Sab commands **read-only** hain. Secrets ki value kabhi print mat karo. AWS par RDS private subnets mein hai (SG sirf task SG se) aur Qdrant/Redis task ke andar hain (task SG sirf ALB se), ECS Exec off – wahan API aur logs use karo.

```
S3
|
PostgreSQL (local)
|
Qdrant (local)
|
Redis (local)
|
Cognito (AWS)
|
ECS / CloudWatch (AWS)
|
AWS par app data: admin token se GET /api/v1/admin/metrics , GET /api/v1/documents , GET /api/v1/admin/audit .
```

32. Real End-to-End Example – "ABC Insurance"

Ye story **currently supported** steps use karti hai. AWS par step 10 aur 20 abhi Bedrock quota (0) ki wajah se fail honge; local mode mein poora flow local AI stub ke saath chalta hai.

#	Kya hota hai	Kaun/kahan
0	Pehla platform operator banao (poore system mein ek hi baar)	<code>./scripts/bootstrap-platform-admin.sh --email ops@...</code>
1	Operator login karta hai, pehle sign-in par apna password set karta hai	Frontend → <code>/auth/login</code> → <code>/auth/new-password</code>
2	Companies → Onboard a company: "ABC Insurance Private Limited", tenant id <code>abc-insurance</code>	<code>POST /api/v1/platform/tenants</code>
3	Invite admin: <code>admin@abc-insurance.example</code> . Cognito one-time password email karta hai – operator wo password kabhi nahi dekhta	<code>POST /platform/tenants/abc-insurance/admins</code>
4	Admin login karta hai, apna password set karta hai, apne users banata hai (HR / Finance / Legal)	<code>POST /api/v1/admin/users</code>
5	Documents page → "HR Maternity Policy.pdf" upload (dept <code>hr</code>)	<code>POST /api/v1/documents</code>
6	File S3: <code>tenant-abc-insurance/<uuid>.pdf</code> (SSE)	<code>storage.put_object</code>
7	Rows: documents pending, ingestion_jobs queued, audit document.uploaded	Postgres (AWS: RDS – ye rows task restart ke baad bhi rehti hain)
8	Background: text extract (pypdf, page numbers ke saath)	<code>extractors.extract_pdf</code>
9	Chunks (512 tokens, 64 overlap) + injection scan	<code>chunker</code> , <code>scan_content</code>
10	Embeddings (Titan v2, batch 25)	Bedrock
11	Qdrant upsert – har point par <code>tenant_id=tenant-abc-insurance</code>	<code>vector.upsert_chunks</code>
12	Document ready, chunk_count set, cache invalidate	Postgres (RDS), Redis
13	(Users step 4 mein hi ban gaye – admin ke Users page se)	<code>POST /api/v1/admin/users</code>
14	User email + password se login karta hai	Frontend → <code>/auth/login</code>
15	Sawaal: "Meri maternity leave policy kya hai?"	Chat page → <code>POST /api/v1/chat</code>
16	JWT verify → <code>tenant_id=tenant-abc-insurance</code> , role user	<code>core/auth.py</code>
17	Validation + injection scan + embed + cache lookup (ABC namespace)	pipeline stages 2–4
18	Qdrant search sirf ABC ke chunks → threshold → rerank	stages 5–10
19	Context envelope → Nova Lite	stages 8–9
20	Citation validation + output guardrail + confidence	stages 11–12
21	Usage + conversation save; answer (language of question) + [S1] → "HR Maternity Policy.pdf, page N"	<code>chat.py</code>
22	Frontend answer, citations, confidence, model, tokens, cost dikhata hai	<code>app/chat/page.tsx</code>

PLATFORM OPERATOR kar sakta: companies banana, unka pehla admin invite karna, registry aur onboarding trail dekhna, company suspend/reactivate karna. **ADMIN kar sakta:** apni company ke users invite/manage karna (role, department, deactivate, password reset), upload, tenant ka koi bhi document delete, chat, workflows, AI metrics, audit, security events. **USER kar sakta:** upload, apna document delete, chat, workflows, feedback. Admin pages nahi. **Koi nahi kar sakta:** doosre tenant ka data dekhna – na admin, na platform operator. Aur `platform_admin` role kisi bhi API se nahi mil sakta.

33. Complete Command Cheat Sheet

Sab commands **repo root** (D:\infy__masters) se, Git Bash mein – jab tak `cd` na likha ho. <...> wali values apni daalo; secrets kabhi command line par mat likho.

Local development

Kaam	Command
Ek-command setup	<code>./scripts/setup-local.sh</code> (PowerShell: <code>.\scripts\setup-local.ps1</code>)
Containers start	<code>docker compose --env-file .env -f infra/docker/docker-compose.yml up -d postgres qdrant redis minio minio-init</code>
Containers status	<code>docker compose --env-file .env -f infra/docker/docker-compose.yml ps</code>
Containers stop (data rahe)	<code>docker compose --env-file .env -f infra/docker/docker-compose.yml stop</code>
Containers + data delete	<code>... down -v</code>
venv activate	<code>cd backend && source .venv/Scripts/activate</code>
Migrations	<code>cd backend && alembic upgrade head</code>
Seed data	<code>cd backend && python -m seeds.seed</code>
Backend run	<code>cd backend && uvicorn app.main:app --reload --port 8000</code>
Frontend run	<code>cd frontend && npm run dev</code>
Login	Browser: seeded email + <code>.env</code> ka <code>DEV_AUTH_PASSWORD</code>
Dev token (user / admin / platform)	<code>cd backend && python -m seeds.dev_token / ... --role admin --user seed-admin-a / ... --role platform_admin</code>
Health	<code>curl localhost:8000/api/v1/health</code> · <code>curl localhost:8000/api/v1/health/ready</code>

Testing

Kaam	Command
Backend lint + format	<code>cd backend && ruff check . && ruff format --check .</code>
Backend sab tests	<code>cd backend && ENVIRONMENT=dev AI_PROVIDER=local DEV_AUTH_ENABLED=true pytest tests/unit tests/integration tests/security tests/evaluation -q</code>
Sirf security	<code>cd backend && pytest tests/security -q</code>
Onboarding journey	<code>cd backend && pytest tests/e2e -q</code> (PostgreSQL chahiye)
Frontend	<code>cd frontend && npm run lint && npm run typecheck && npm run build</code> (dev server band karke)
Cost-guard Lambda	<code>backend/.venv/Scripts/python.exe -m pytest infra/terraform/modules/cost-guard/lambda/tests -q</code>

AWS lifecycle

Kaam	Command
Kharcha + kya chal raha hai	<code>./scripts/cost-check.sh</code>
Pehla platform operator	<code>./scripts/bootstrap-platform-admin.sh --email ops@... [--dry-run]</code>
Backend deploy	<code>./scripts/deploy.sh</code>
Verify	<code>./scripts/verify.sh</code>
Seed (AWS)	<code>./scripts/seed.sh</code>
E2E	<code>./scripts/test-e2e.sh</code>
Backend rollback	<code>CodeDeploy auto-rollback · ya git revert <BAD_SHA> + ./scripts/deploy.sh</code> (Chapter 22.5)
<code>rollback.sh</code>	Sirf verify + report – traffic shift implemented nahi
Frontend deploy / verify	<code>git push origin aws-deployment · ./scripts/deploy-frontend.sh</code>
Frontend rollback	<code>./scripts/rollback-frontend.sh</code> (ya <code>--to <COMMIT_SHA></code>)
Destroy backend	<code>./scripts/destroy.sh → DESTROY ekba-dev</code>
Identity check	<code>aws sts get-caller-identity --query Account --output text</code>
Mera IP	<code>curl -s https://checkip.amazonaws.com</code>

Diagnose (read-only)

Kaam	Command
ECS service	<code>aws ecs describe-services --cluster ekba-dev --services ekba-dev --query 'services[0].[status,runningCount,desiredCount]'</code>
ECS events	<code>aws ecs describe-services --cluster ekba-dev --services ekba-dev --query 'services[0].events[:5].[createdAt,message]' --output text</code>
Backend logs	<code>MSYS_NO_PATHCONV=1 aws logs tail /ekba/dev/service --since 30m</code>
Ek request trace	<code>MSYS_NO_PATHCONV=1 aws logs tail /ekba/dev/service --since 1h --filter-pattern '"<CORRELATION_ID>"'</code>
Target health	<code>aws elbv2 describe-target-health --target-group-arn <TG_ARN></code>
CodeDeploy latest	<code>aws deploy list-deployments --application-name ekba-dev --deployment-group-name ekba-dev-dg --query 'deployments[0]' --output text</code>
CodeDeploy status	<code>aws deploy get-deployment --deployment-id <DEPLOYMENT_ID> --query 'deploymentInfo.[status,errorInformation]'</code>
Amplify jobs	<code>aws amplify list-jobs --app-id <APP_ID> --branch-name aws-deployment --query 'jobSummaries[:3].[jobId,status]'</code>
CloudFront → API	<code>curl -s https://<CLOUDFRONT_DOMAIN>/api/v1/health</code>
Cost-guard logs	<code>MSYS_NO_PATHCONV=1 aws logs tail /aws/lambda/ekba-dev-cost-guard --since 1h</code>
RDS status	<code>aws rds describe-db-instances --db-instance-identifier ekba-dev-postgres --query 'DBInstances[0].[DBInstanceStatus,PubliclyAccessible,StorageEncrypted]' --output text</code>
RDS snapshots	<code>aws rds describe-db-snapshots --db-instance-identifier ekba-dev-postgres --snapshot-type manual --query 'DBSnapshots[].[DBSnapshotIdentifier,Status]' --output text</code>
Database reachable? (task se)	<code>curl -s https://<CLOUDFRONT_DOMAIN>/api/v1/health/ready - postgres component dekho</code>
Terraform outputs	<code>cd infra/terraform/envs/<env> && terraform output</code>

Git

Kaam	Command
Status	<code>git status</code>
Kya badla	<code>git diff / git diff --cached --check</code>
Commit (files naam se)	<code>git add <files> && git commit -m "<message>"</code>
Frontend release	<code>git push origin aws-deployment</code>
Branch compare	<code>git rev-list --count origin/main..aws-deployment</code>

34. Interview Preparation

34.1 Two-minute version

"Maine ek **multi-tenant enterprise RAG platform** banaya hai. Companies apne documents – PDF, Word, Excel, images, audio – upload karti hain, aur employees natural language mein sawaal poochte hain; jawab **citations aur confidence score** ke saath aata hai.

Backend **FastAPI** hai. Upload par file validate hoti hai (extension + magic bytes), S3 mein encrypted jaati hai, text extract hota hai, 512-token chunks bante hain, **Titan v2** se 1024-dim embeddings **Qdrant** mein jaati hain. Query par ek **14-stage pipeline** chalti hai: injection scan, semantic cache, tenant-filtered vector search, relevance threshold, BM25 rerank, **Amazon Nova** se answer, phir citation validation aur output guardrail.

Sabse important cheez **tenant isolation** hai – **tenant_id** sirf verified **Cognito JWT** se aata hai aur har DB query, har Qdrant search, har cache key aur har S3 prefix mein lagta hai.

AWS par **ECS Fargate** chalta hai, **CodeDeploy blue/green** ke saath auto-rollback; frontend **Amplify** par static export, API **CloudFront** ke through HTTPS. Sab **Terraform** mein – chaar states. Budget \$20 tha, isliye ephemeral design: NAT aur ElastiCache nahi; Qdrant aur Redis containers hain, aur PostgreSQL ek chhota **RDS (db.t4g.micro)** jo stack ke saath hi banta-mitta hai – uska data snapshot ke through agle session mein wapaa aa jaata hai. Ek **cost-guard Lambda** hai jo budget cross hone ya 8 ghante baad stack band kar sakta hai."

34.2 Five-minute version – structure

1. **Problem** (30 s): company knowledge scattered; generic chatbots hallucinate aur tenants ka data mix kar sakte hain.
2. **Architecture** (60 s): Browser → Amplify (static Next.js) · API → CloudFront → ALB → ECS task (api + qdrant + redis) → RDS PostgreSQL (private subnets), Bedrock, S3, Cognito, Transcribe.
3. **Ingestion** (45 s): validation, S3 key <tenant>/<uuid>, extractors, chunker, per-chunk injection scan (retrieval poisoning), batch embeddings, Qdrant with tenant_id payload index.
4. **RAG + guardrails** (75 s): 14 stages; context envelope (documents = data, not instructions); citations [Sn] validate; confidence formula; semantic cache per tenant (cosine ≥ 0.95); daily cost ceiling per user.
5. **Security** (45 s): JWT RS256 via JWKS; roles; rate limits; audit events; security headers; non-root container; secrets in Secrets Manager.
6. **DevOps** (45 s): Terraform 4 states; blue/green with 5xx alarm rollback; CI (lint, tests, gitleaks, Trivy, checkov, terraform validate).
7. **Honest limits** (30 s): Bedrock quota 0 is live par; relational data ab RDS mein bachta hai par Qdrant vectors abhi bhi task ke saath jaate hain; no self-service signup; hosted-UI login nahi; CI main par.

34.3 Likely questions – short answers

Question	Answer
RAG kya hai?	Model ko sawaal ke saath relevant document chunks dena taaki jawab un par based ho, training data par nahi
Qdrant kyun?	Payload filtering (tenant) + open source + container mein chal jaata hai (\$0)
Tenant isolation kaise?	Tenant token se; har layer (DB, Qdrant must filter, Redis key, S3 prefix) mein; cross-tenant attempt par critical audit; tests/security
Hallucination kaise kam?	Relevance threshold, "not found" refusal, citation validation, confidence score, strict prompt
Prompt injection?	Input scan (10 families), chunk scan at ingestion, envelope, output guardrail
Semantic cache?	Question embedding ka cosine ≥ 0.95 match same tenant mein; cache hit par bhi auth + guardrail
Blue/green?	Naya task set test listener (:8080) par; healthy hone par traffic shift; 5 minute rollback window; 5xx alarm par auto-rollback
Database kaunsa?	AWS par RDS PostgreSQL (db.t4g.micro) , private, encrypted, password Secrets Manager se write-only argument ke through). Pehle sidecar container tha – data har task restart par chala jaata tha
Database sirf ek chhota instance kyun?	\$20 ceiling. Single-AZ db.t4g.micro ~\$14/month 24/7, isliye wo ephemeral stack ke saath hi rehta hai aur destroy.sh snapshot le leta hai. Production mein Multi-AZ + read replica + ElastiCache
Scale kaise?	Stateless API ke multiple tasks; managed DB/Qdrant Cloud; ingestion ko SQS worker mein nikalna
Biggest challenge?	CodeDeploy Free plan block; IP change par health fail; Amplify monorepo headers – har ek ka root cause dhoondha
Multi-tenancy kaise onboard hoti hai?	Teen-level hierarchy: platform_admin company banata hai aur uska pehla admin invite karta hai; wo admin sirf apni company ke users manage karta hai. Har level sirf apne se ek neeche wale ko bana sakta hai
Tenant admin ko platform admin banne se kaise roka?	Chaar layers: role schema mein hi valid value nahi (422), TENANT_ASSIGNABLE_ROLES mein nahi, server-side assert_role_assignable refuse karta hai, aur token verify par platform role sirf platform tenant ke saath valid hai. Koi API platform role de hi nahi sakti
Platform admin tenant isolation ko tod nahi deta?	Nahi – wo apne reserved platform tenant mein rehta hai (jisme content nahi hai), aur uska ekmatra cross-tenant access control_plane.py hai: sirf registry data (naam, id, seat counts), koi document/chat/metric nahi, aur har mutation audited
Password kaise handle karte ho?	Karte hi nahi. Cognito one-time password email karta hai, user pehle sign-in par apna set karta hai. App password set/store/log/return kabhi nahi karta, aur task role ke paas AdminSetUserPassword hi nahi hai
Login ko brute force se kaise bachaya?	Per-account 10 attempts/min (IP par nahi – CloudFront ke peeche IP shared ya client-supplied hoti hai), generic failure message, aur forgot-password hamesha 202
Kya improve karoge?	SES se branded invitation emails, MFA (pool OFF par hai), SCIM/SSO se bulk user provisioning, SQS ingestion worker, WAF + custom domain, main-based CI/CD

Interview mein **wahi bolo jo implemented hai**. "Async workers (separate queue)", "LangSmith tracing", "Cognito hosted UI", "MFA" – ye abhi **implemented nahi** hain; poochha jaye to "planned" bolo. Login, refresh, aur tenant/user admin APIs **implemented** hain.

35. NEVER DO THIS

In mein se koi bhi galti data leak, kharcha ya toota hua deploy de sakti hai.

1. Secrets commit mat karo - `.env`, `*.tfvars`, `*.tfstate`, `tfplan`, `backend.hcl`, `keys`, `tokens`.
2. `git add .` mat karo - `docs/reports/2026*.md` (account ID + live URLs) andar aa jaati hain. Files naam se add karo.
3. `terraform destroy` seedha mat chalao - sirf `./scripts/destroy.sh`.
4. `Baseline / cost-guard / frontend / state backend destroy` mat karo - secrets, ECR images, Cognito users, budget, URLs sab jayenge.
5. `Destroy/replace` wala plan bina samjhe `apply` mat karo - khaaskar Amplify app replace (naya URL, GitHub reconnect).
6. `dry_run = false` bina soche mat karo - agli hourly run 8 ghante purana stack band karegi.
7. AWS Budget ko hard cap mat samjho.
8. Frontend se `tenant_id` bhejne wala endpoint mat banao - tenant sirf JWT se.
9. User ka `custom:tenant_id` galti se mat badlo - wo user doosri company ka data dekhne lagega.
10. `Guardrail stage skip` mat karo "speed ke liye", aur test delete karke build green mat karo.
11. `NEXT_PUBLIC_*` mein secret mat daalo - browser mein sab public hai.
12. Real JWT online decoder mein paste mat karo.
13. ECS ko haath se `scale / console` se edit mat karo - Terraform drift; ALB ka kharcha chalta rehta hai.
14. Dev server chalte hue `npm run build` mat chalao.
15. `docker compose down -v` bina soche mat chalao - local DB, Qdrant, MinIO data delete.
16. Embedding model badal kar same Qdrant collection use mat karo - nayi collection + re-index.
17. Stack chalu chhod kar mat jao - session ke end mein `./scripts/destroy.sh`.
18. Kisi aur project ke AWS resources ko touch mat karo - `ownership = state + ekba-* naam + ProjectCode=ekba` tag.

36. Glossary

Term	Simple matlab (is project mein)
RAG	Retrieval-Augmented Generation – pehle documents se relevant hisse dhoondo, phir model unke basis par jawab de
LLM	Large Language Model – yahan Amazon Nova Lite / Micro
Embedding	Text ka number-vector (1024 numbers, Titan v2) – similar meaning = paas-paas vectors
Vector DB	Vectors store + similarity search – yahan Qdrant
Chunk	Document ka chhota tukda (~512 tokens) jo embed hota hai
Token	Text ki unit (~4 characters); model ka kharcha tokens par
Cosine similarity	Do vectors kitne same direction mein – 1 = same meaning
Rerank	Vector results ko dobara order karna (yahan BM25 lexical mix)
Semantic cache	Milte-julte sawaal ka pichla jawab reuse (same tenant)
Tenant	Ek company; uska data baaki se alag
JWT	Signed token jisme user ki claims (tenant, role) hoti hain
ID token vs access token	Cognito ke do token; backend ko ID token chahiye (usme aud + custom claims)
JWKS	Cognito ki public keys jinse token signature verify hota hai
Cognito	AWS user directory + login
Guardrail	Safety check stage (input, retrieval, output)
Prompt injection	Text jo model ko instructions badalne ke liye bahkaaye
Citation [Sn]	Jawab ka woh hissa kis source chunk se aaya
ECS / Fargate	AWS container service / bina server manage kiye containers
Task definition	Container recipe (image, CPU, memory, env, secrets)
ALB	Application Load Balancer – traffic ko healthy tasks par bhejta
Target group	ALB ke peeche tasks ki list (blue / green)
CodeDeploy blue/green	Naya version saath mein chalao, check karo, traffic shift karo
CloudFront	AWS CDN – yahan API ko HTTPS dene ke liye
Amplify	Static frontend hosting + GitHub se auto-build
Static export	Next.js ko plain HTML/JS files mein build karna (NEXT_OUTPUT_MODE=export)
CORS	Browser rule – kaunsi website API call kar sakti hai
Terraform state	Terraform ki yaaddasht – kaunsa resource usne banaya
Plan / apply	Kya badlega dikhana / sach mein badalna
Drift	Asli AWS aur Terraform state mein farak (haath se change karne se)
Ephemeral	Temporary – destroy + recreate by design
DRY_RUN	Cost guard sirf batata hai, kuch band nahi karta
Correlation ID	Ek request ka unique ID, har log line mein
Alembic	Database schema migrations tool
Idempotent	Dobara chalane par same result, koi nuksaan nahi

37. Final Quick Reference

Cheez	Value
Frontend (live)	https://aws-deployment.d39pgpamq0p0n8.amplifyapp.com
API (live, jab backend deployed ho)	https://d2jw2wchz5oekh.cloudfront.net/api/v1
Local	Frontend http://localhost:3000 · API http://localhost:8000 · Docs /docs · Qdrant :6333/dashboard · MinIO :9001
Region / states	us-west-2 · baseline, cost-guard, dev, frontend
Release branch	aws-deployment (Amplify) · CI main par
Login	Email + password (dono jagah). Local password: .env ka DEV_AUTH_PASSWORD
Token	Cognito ID token, /auth/login se – frontend khud manage karta hai. Debug ke liye python -m seeds.dev_token (local only)
Pehla operator	./scripts/bootstrap-platform-admin.sh --email ... (iski koi API nahi hai)
Cost guard	DRY_RUN=true · \$18 gross shutdown · card > \$0.01 · 8 ghante

Roz ka flow

Problem aaye to – pehle ye 5

1. ./scripts/cost-check.sh – stack chal bhi raha hai?
2. curl -s https://checkip.amazonaws.com vs allowed_cidrs – IP badla?
3. MSYS_NO_PATHCONV=1 aws logs tail /ekba/dev/service --since 30m
4. ECS service events (describe-services ... events[:5])
5. Amplify job logs / CodeDeploy lifecycle events

Files jo sabse zyada kaam aayengi

backend/app/services/rag/pipeline.py (RAG) · backend/app/core/auth.py (login) · backend/app/core/config.py (settings) · frontend/src/lib/api.ts (API client) · infra/terraform/envs/dev/terraform.tfvars (IP, sizing) · infra/terraform/envs/cost-guard/terraform.tfvars (DRY_RUN) · RUNBOOK.md · scripts/

Mujhe ye karna hai – kaunsa chapter?

Kaam	Chapter	Kaam	Chapter
Project local start	18	Backend change / deploy	20.1 / 19.2
Company onboard karna	7.3	Logs dekhna	27
Users banana / use karna	7, 6	Login / password troubleshoot	28.2 (Auth)
Role hierarchy samajhna	2.2, 8.1	API troubleshoot	28.2
Document add karna	12, 32	Cognito / ID token	6.6, 28.2
Data kahan jaata hai	17	Qdrant troubleshoot	15, 28.2
Sawaal poochna	13.1, 32	Frontend / backend stop	18.5, 22, 22.5
RAG flow samajhna	13	AWS band karna	22
Data inspect karna	31	Redeploy	22.4
Frontend change	20.2	Rollback	22.5
Frontend deploy	19.1, 22.2	Safe destroy	22.3, 24.3
AWS kharcha	30	Interview	34

Yaad rakho: Local = \$0, sab kuch pehle local mein. AWS = demo ke liye, aur demo ke baad destroy.sh .